Data Engineering

MG-GY 8441

Describing Data

- Agenda
  - Overview
  - Data Types
  - Numerical Summaries
  - Visualizations
- References
  - Han, Kamber, Pei, *Data Mining: Concepts and Techniques*
    - Chapter 1
    - Chapter 2.1 - 2.3

Overview

Overview

- Why Data Mining?
- What Is Data Mining?
- What Kind of Data Can Be Mined?
- What Technology Are Used?
- What Are the Applications?
- What Are the Challenges?

Why Data Mining?

Classification

- Classify credit applicants as low, medium, high risk

Estimation

- Estimate the **click-through-rate** of an advertisement

Prediction

- Predict which customers will leave within six months

Why Data Mining?

Example from Marketing

- You are in a meeting with your boss and a large publisher where you are negotiating to buy some advertising on their website.
- The publisher tells you the cost per thousand views of your advertisement (CPV) is $10.
- Given your goal of collecting email addresses for potential new customers, you need to know the maximum CPV you can afford to effectively negotiate with the publisher.

Why Data Mining?

Example from Marketing

- You estimate that the click-through rate (CTR) of your advertisement has been around 1%.
- Your conversion rate (CR) has been averaging 10% in terms of email sign-ups.
- If you can afford to pay $5 per email you acquire as your cost per acquisition (CPA), is $10 CPV a good price from the publisher?

Why Data Mining?

Example from Marketing:

- CPV = $10, CTR = 1%, CR = 10%, CPA goal = $5
  - CPA = CPV / ((CTR*1000) * CR)
  - CPA = $10 / ((0.01 * 1000) * 0.1)
  - CPA = $10 / (10 * 0.1)
  - CPA = $10 / 1
- So the cost per acquisition goal would need to be doubled to match the publisher's price of $10 for cost per thousand views

What Is Data Mining?

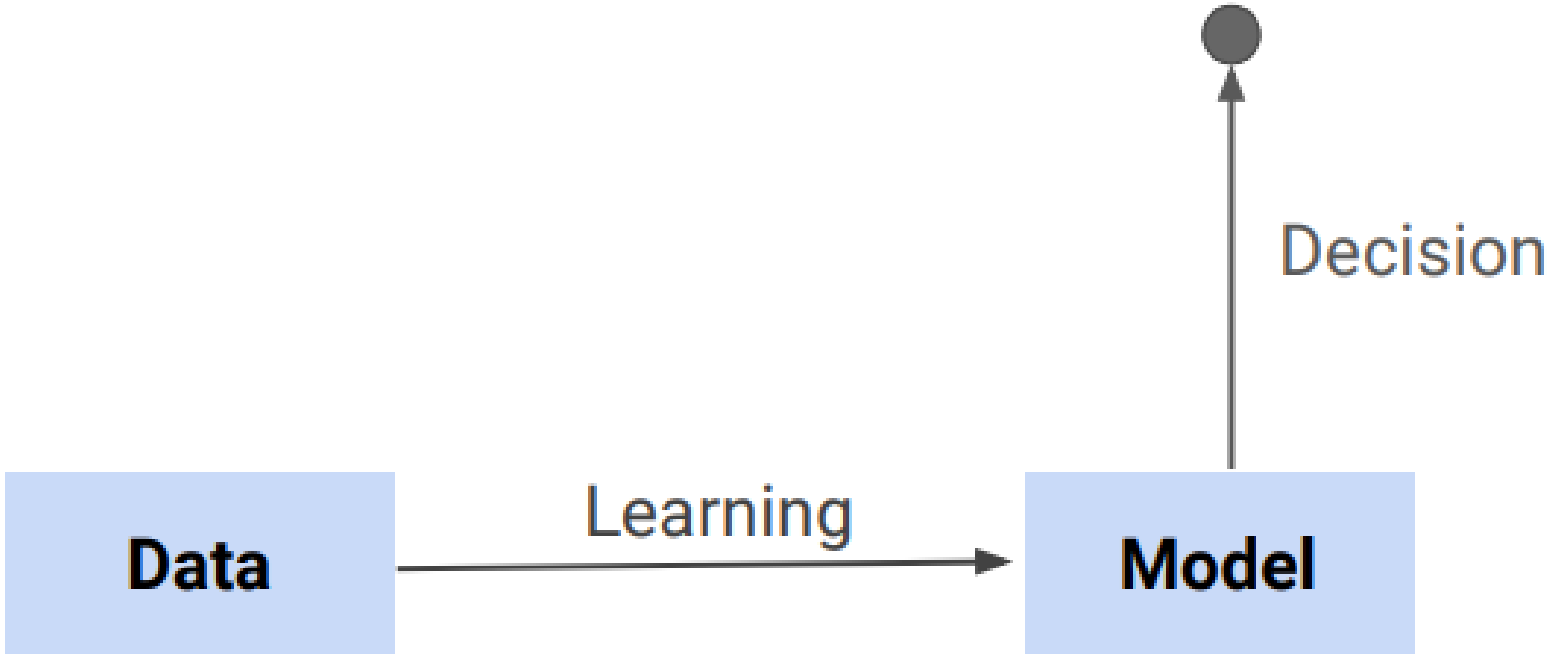

What Is Data Mining?

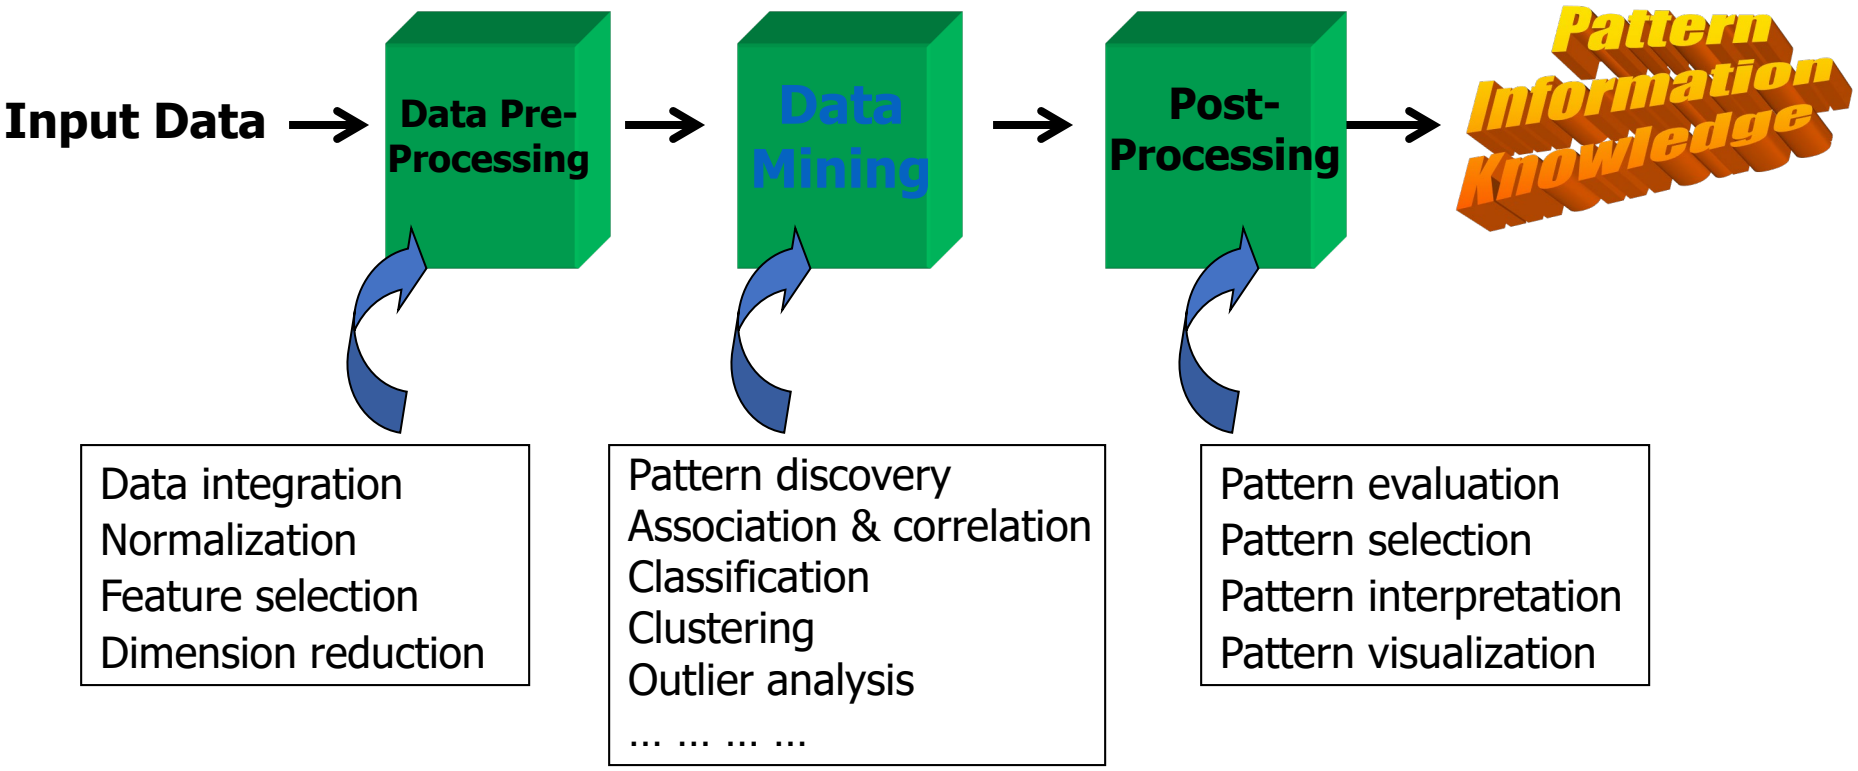

What Is Data Mining?

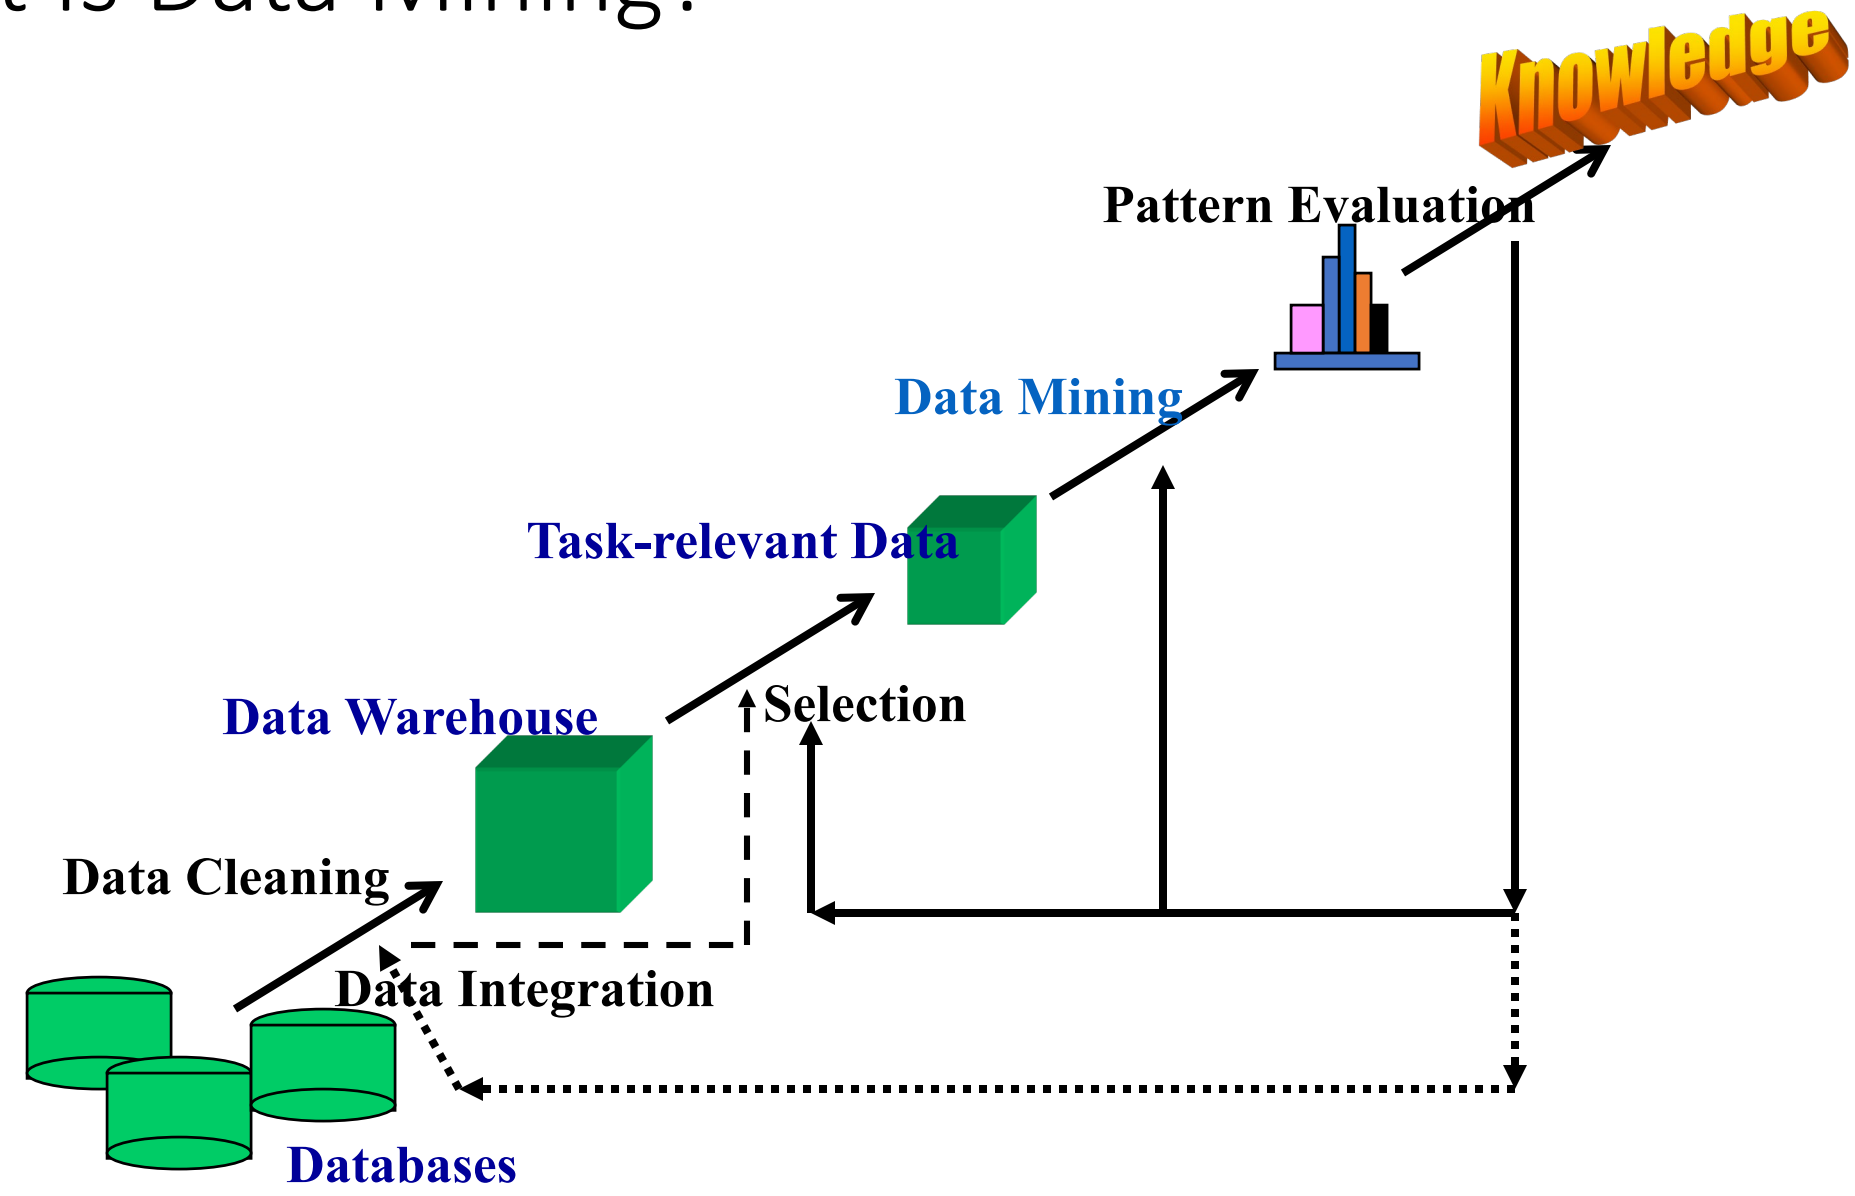

What Is Data Mining?

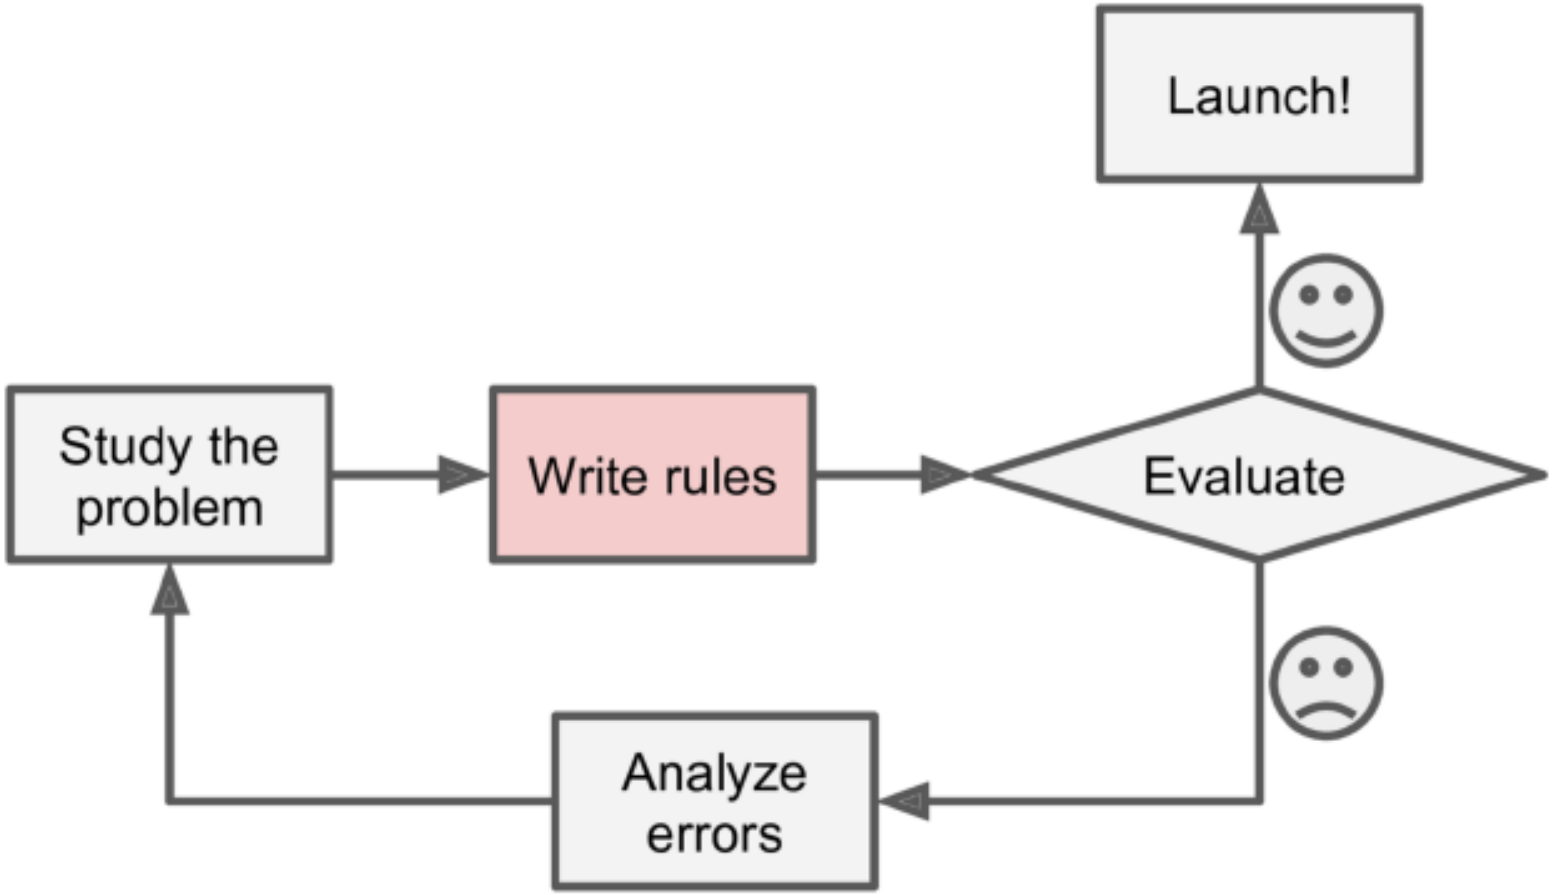

What Is Data Mining?

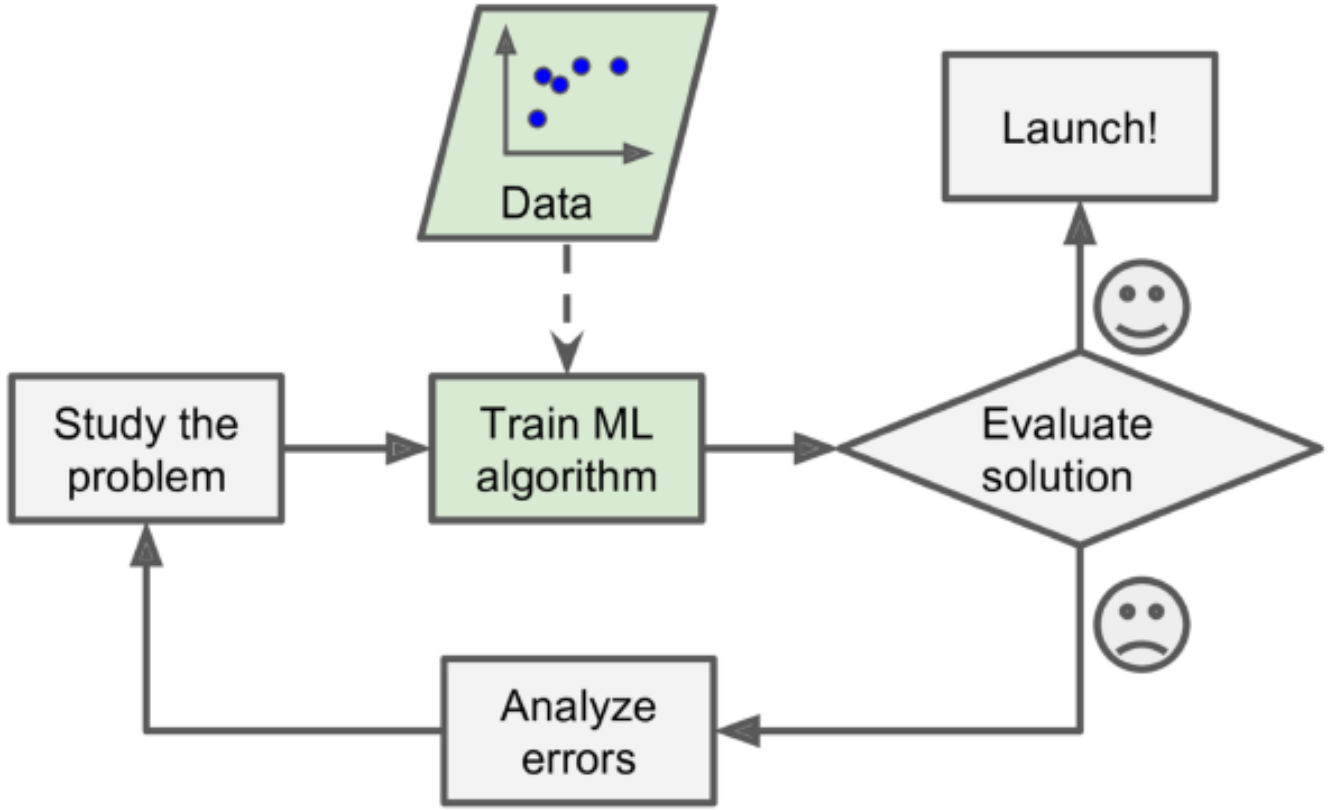

What Is Data Mining?

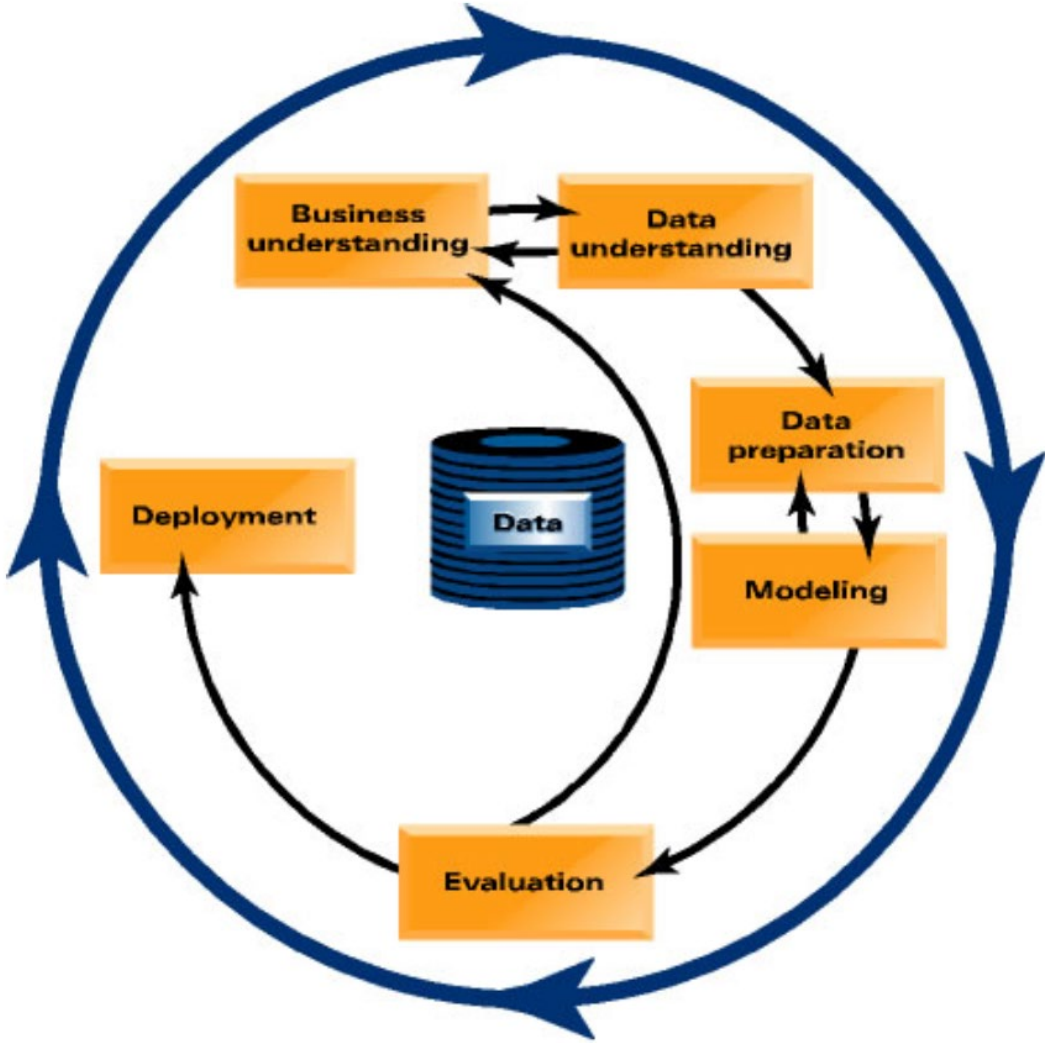

What Kind of Data Can Be Mined?

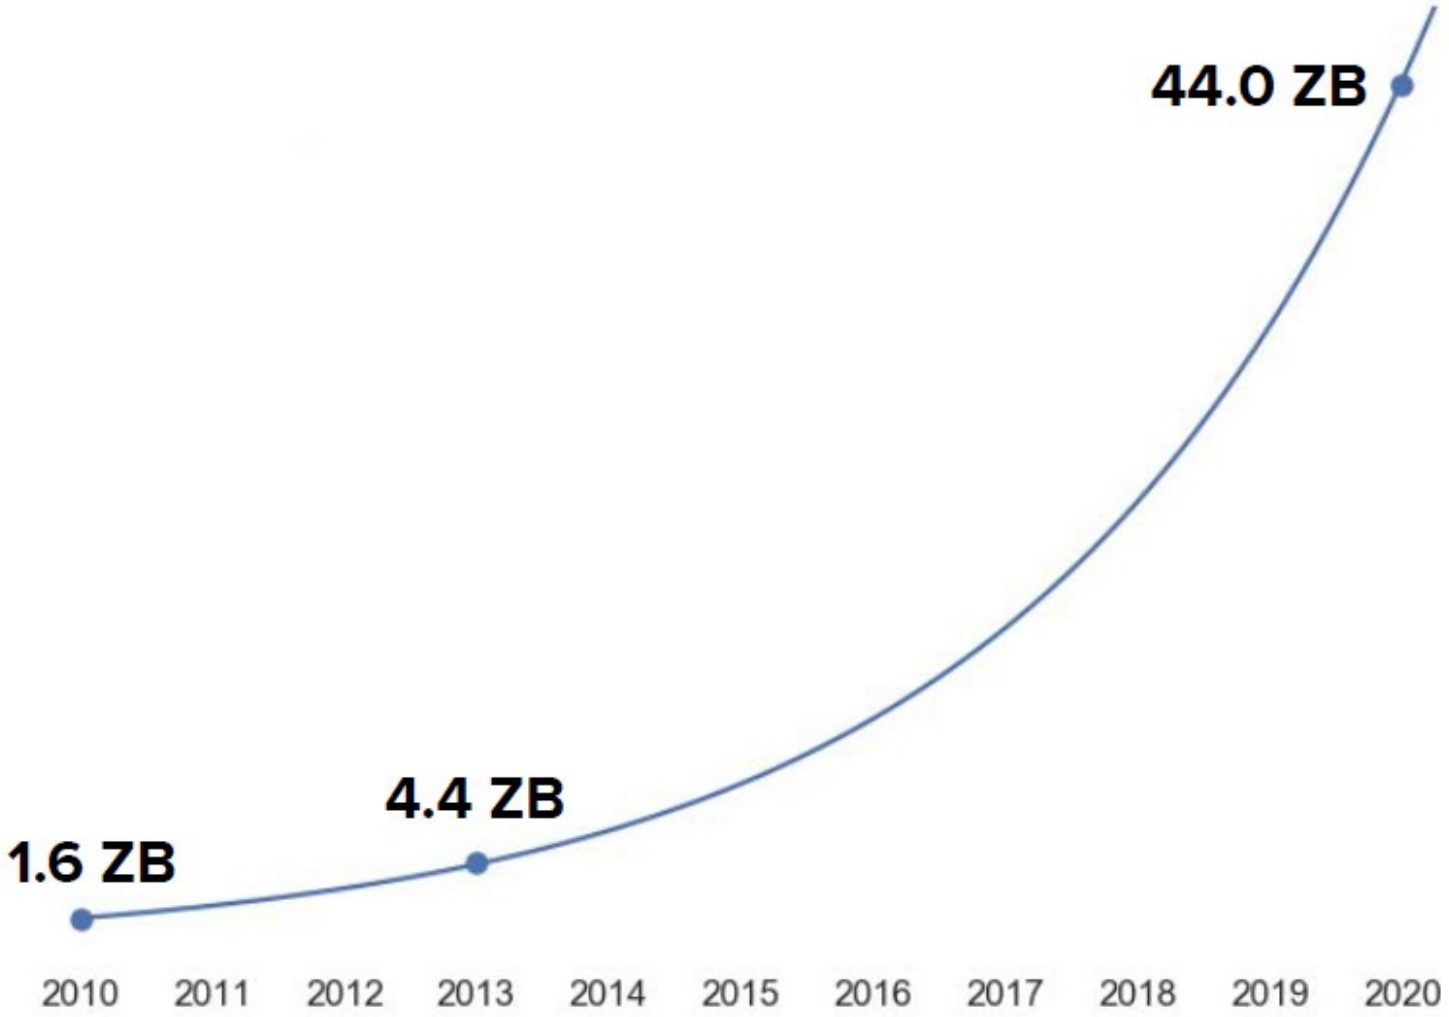

What Kind of Data Can Be Mined?

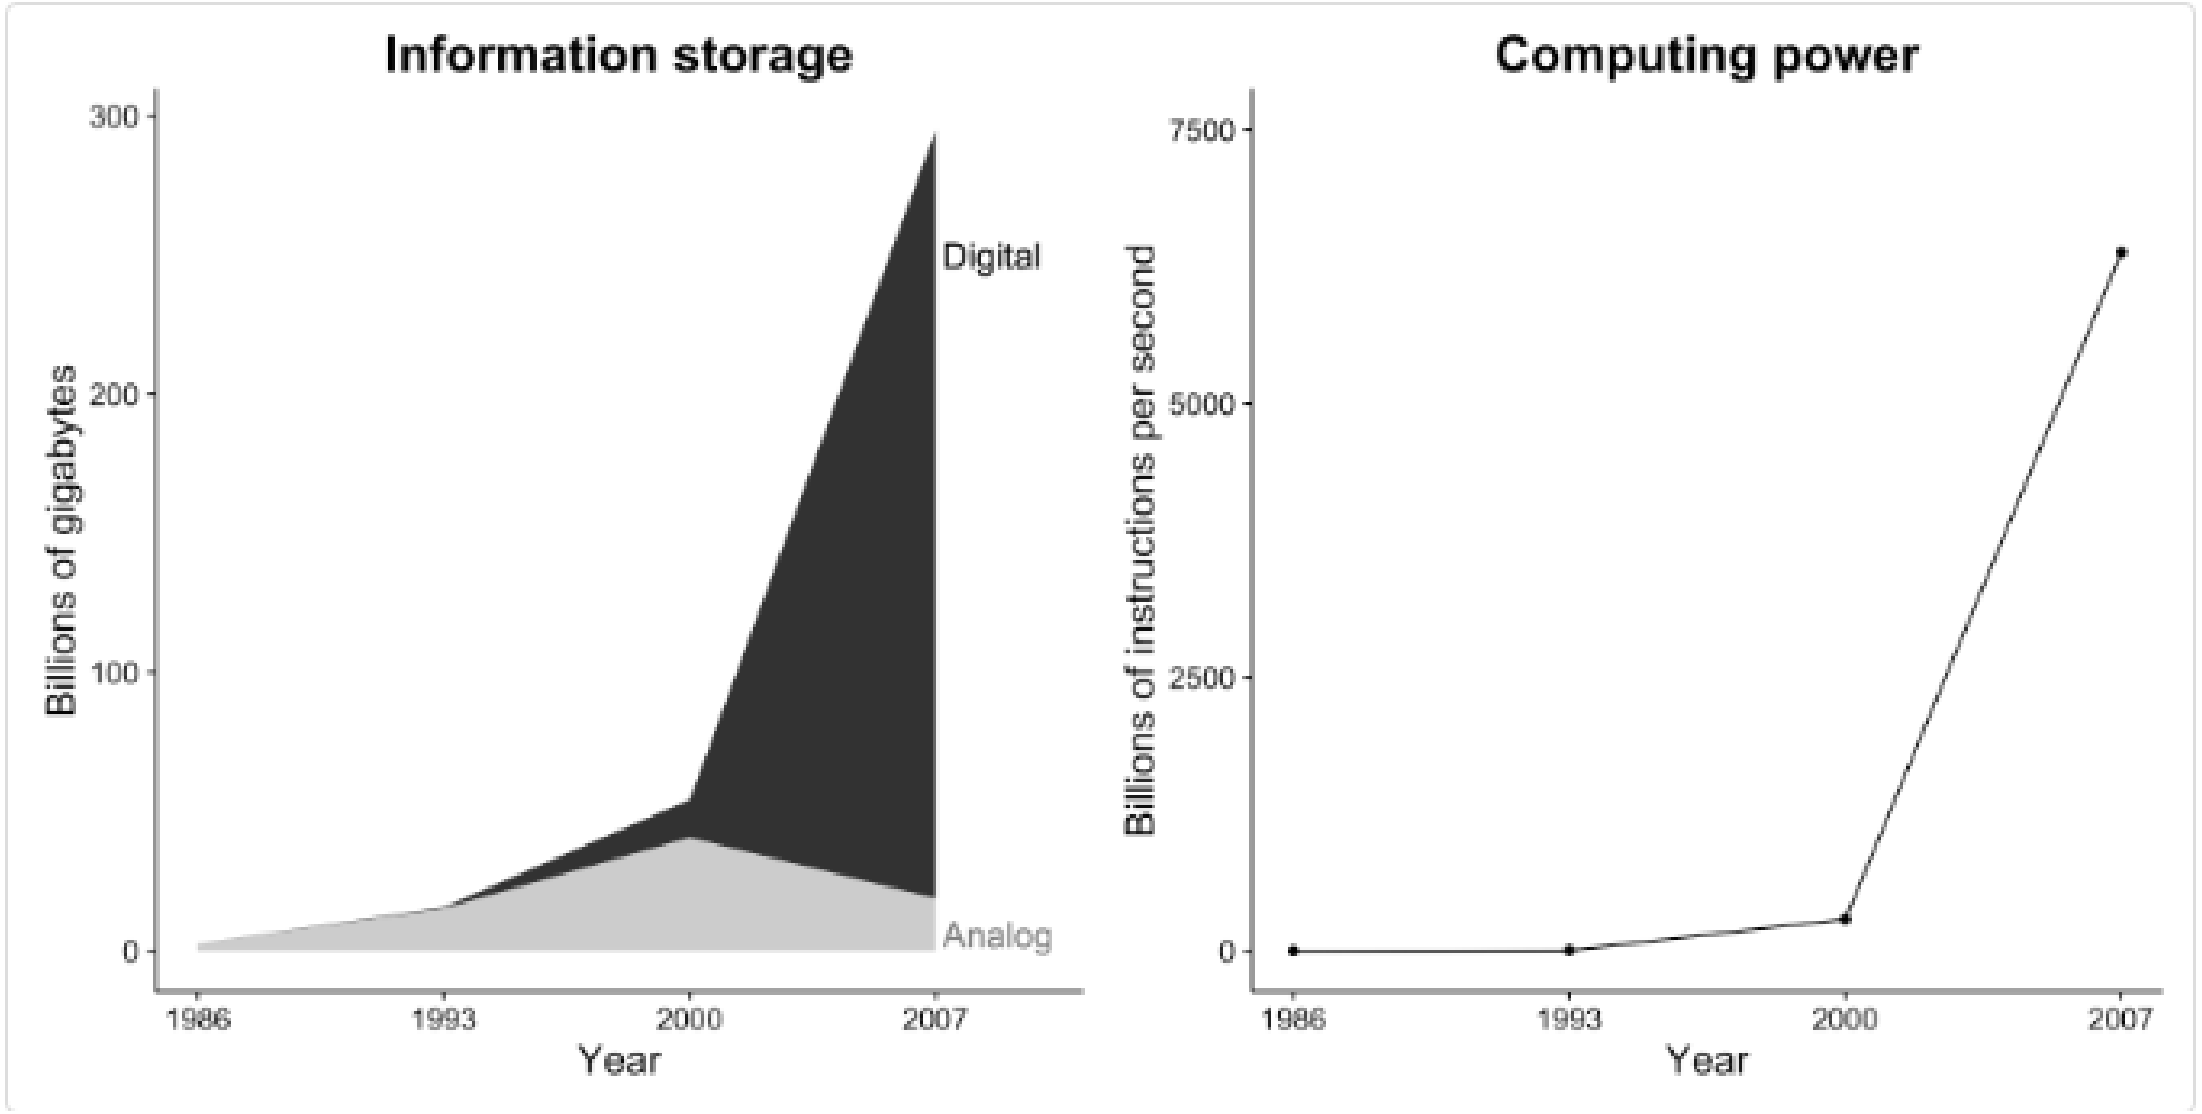

What Technology Are Used?

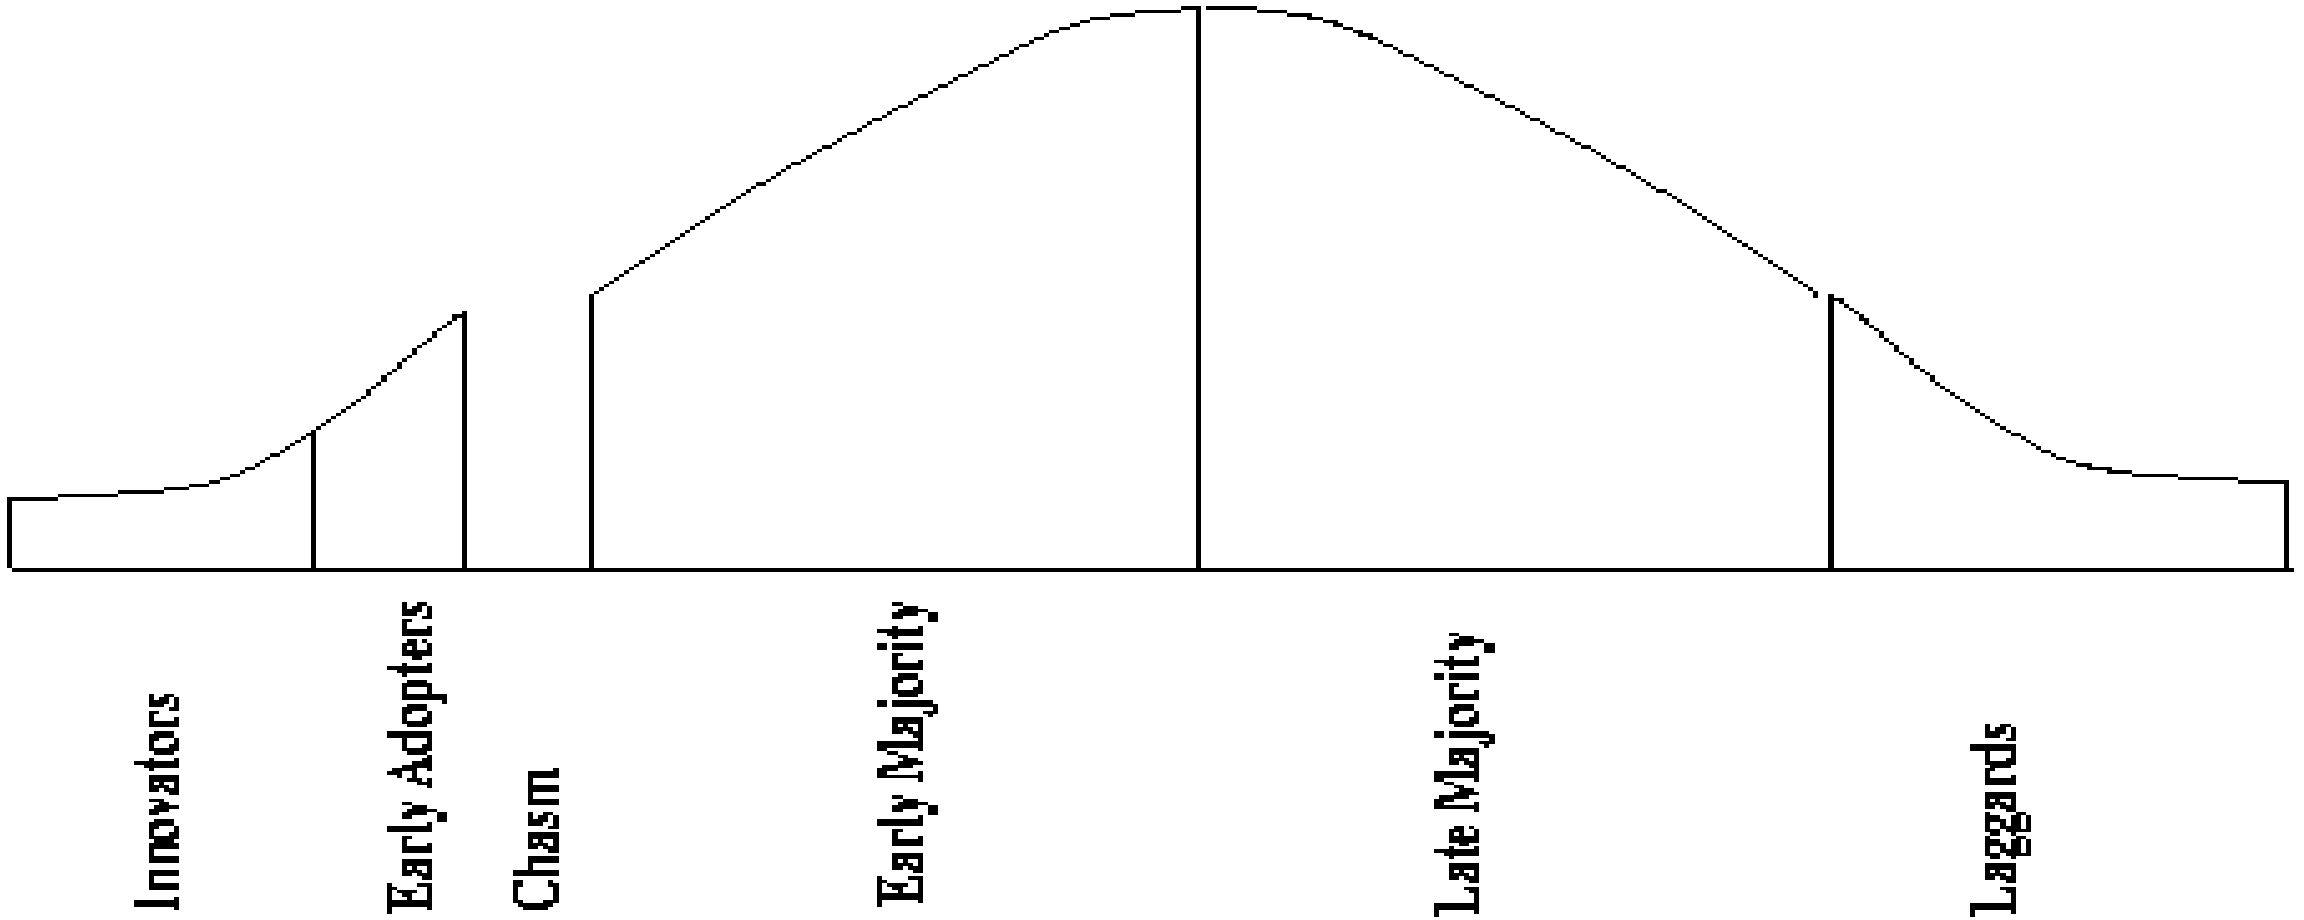

What Technology Are Used?

What Are the Applications?

- Market Basket Analysis
  - Identify what products are likely to be bought together
- Entity Resolution
  - Disambiguate records by linking various data sources
- Market segmentation
  - Identify common characteristics of customers who buy same products
- Collaborative Filtering
  - Recommend products to customers based on preferences

What Are the Applications?

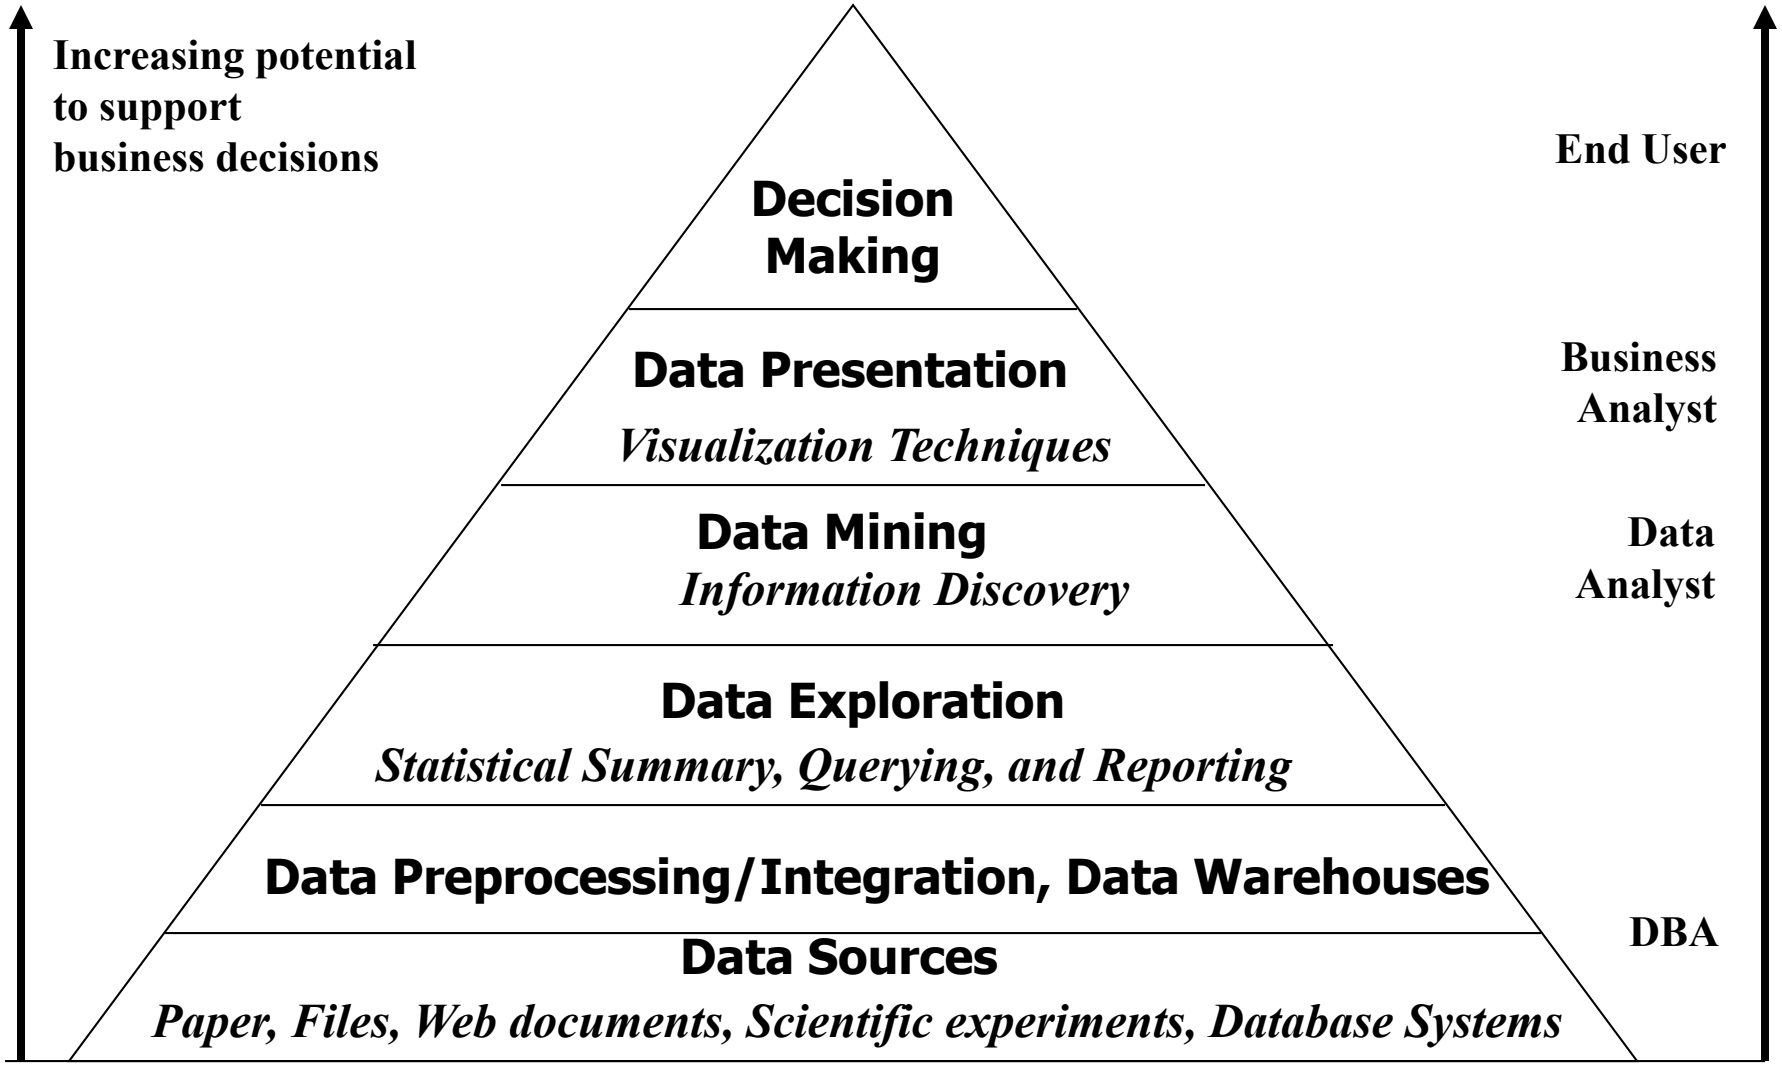

What Are the Challenges?

- Privacy
  - Right to be unknown or forgotten
- Transparency
  - Redistribution of data
- Accountability
  - Oversight of companies and government agencies
- Fairness
  - Social impact of data driven decision making

Data Types

Data Types

- Categorical
  - Ordinal
  - Nominative
- Numerical
  - Continuous
  - Discrete

Properties of Data

Volume	The quantity of data
Velocity	Speed at which data is collected
Variety	Data may be structured or heterogeneous
***Veracity**	Data can be noisy, incomplete, or wrong

Properties of Data

- Dimensionality
  - High dimensional data needs to transform to low dimension data for processing
- Sparsity
  - If many entries lack information, then we need to compress the data to improve storage and computation
- Resolution
  - The granularity of a dataset refers to the level of detail. Sometimes we need to adjust the granularity by grouping records.
- Distribution
  - Aggregations of numbers allow use to track trends in data like dispersion around a common center

Data Object

- Datasets are made up of data objects. A **data object** represents an entity amongst the records.
- Example:
  - sales database: customers, store items, sales…
- Data objects are described by **attributes**.
  - We tend to organize data into tables with rows and columns
    - rows -> data objects
    - columns ->attributes.

Qualitative Data

- **Nominal:** categories, states, or “names of things”
  - Example: marital status, occupation, zip codes
- **Ordinal**
  - Values have a meaningful order (ranking) but magnitude between successive values is not known.
  - Example: *small, medium, large*}

Quantitative Data

- **Discrete Attribute**
  - We can count the distinct numbers.
    - Commonly the numbers have integer values like 0,1,2,3,..
  - Binary attributes are a special case of discrete attributes containing the value 0 or 1
- **Continuous Attribute**
  - We cannot count the range of numbers
    - Commonly the numbers have floating point values containing fractions with many digits

Quantitative Data

- **Interval**
  - Range of numbers differing by increments
  - Example: 1,2,3,4,5 contains numbers differing by increment of 1
- **Ratio**
  - Range of numbers differing by factor
  - Example:
    - 10^1, 10^2 , 10^3 , 10^4 , 10^5 contains numbers differing by factor of 10
    - Factors of 10 are called order of magnitude. This type of range is a logarithmic scale. It can be use for working with large ranges of values in visualizations

Numerical Summaries

Numerical Summaries

- Measuring Central Tendencies
  - Mean
  - Median
  - Mode
- Ranking Numbers
  - Quantiles

Measuring the Central Tendency

- Mean
  - Measures the average of a collection of numbers.
  - If we want to put different emphasis on the numbers, then we can use weighted mean
  - If we want to remove outliers, then we could discard or round some values

$$\overline{x} = \frac{\sum_{i=1}^{n} w_i x_i}{\sum_{i=1}^{n} w_i}$$

Measuring the Central Tendency

- Median:
  - We can rank the numbers from smallest to largest.
  - The median is the smallest number such that at least half of the other numbers of lesser value

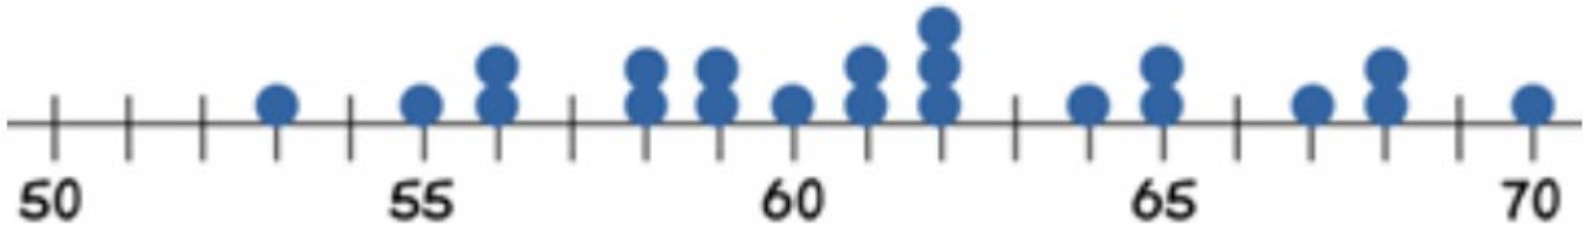

Measuring the Central Tendency

Mode

- Value that occurs most frequently in the data
- Unimodal, bimodal, trimodal,…

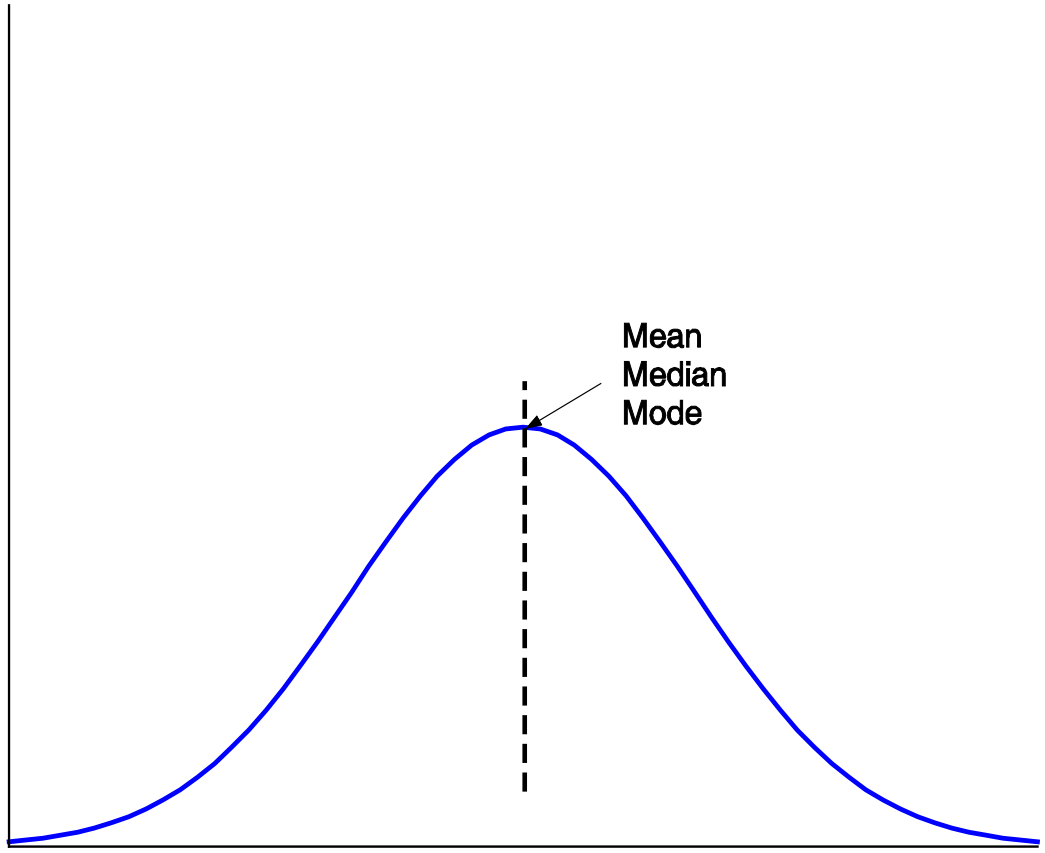

Symmetric vs. Skewed Data

- Median, mean and mode of symmetric, positively and negatively skewed data

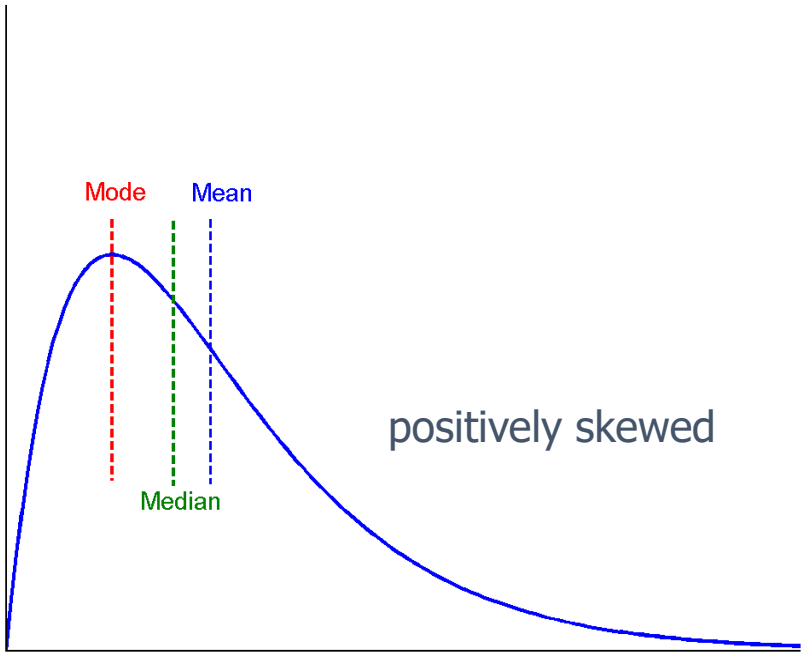

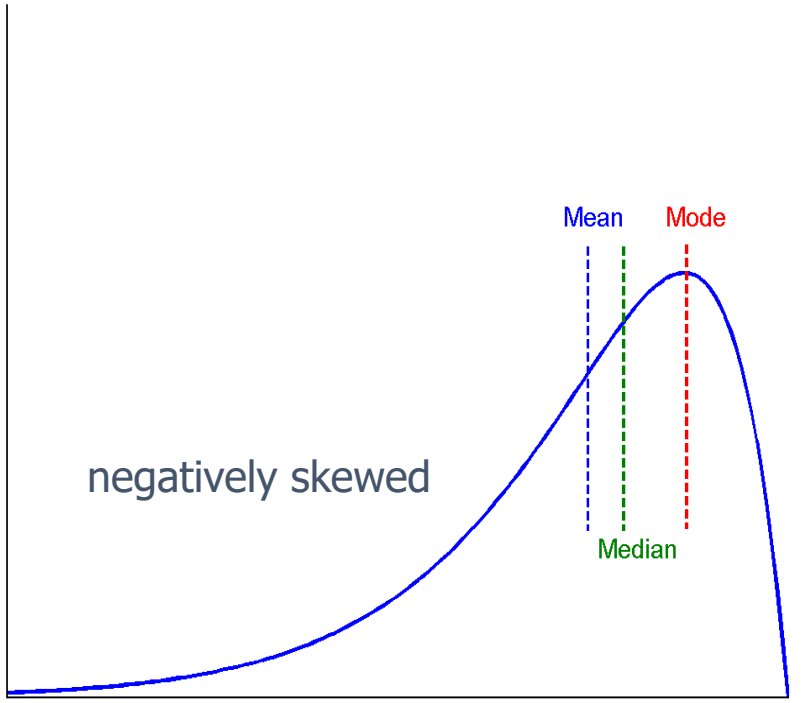

Dispersion of Data

- Quantile
  - Quantiles generalize median. Instead of splitting the data in half, we can split into any fractions.
    - **Quartiles**: Q_1 (25^{th} percentile), Q_3 (75^{th} percentile)
  - **Five number summary**: minimum, Q_1, median, Q_3, maximum
  - **Inter-quartile range**: IQR = $Q_3 - Q_1$
    - **Outlier**: usually, a value higher/lower than 1.5 x IQR

Dispersion of Data

- **Variance** is average square distance around the mean

$$s^2 = \frac{1}{n-1}\sum_{i=1}^{n}(x_i - \bar{x})^2 = \frac{1}{n-1}[\sum_{i=1}^{n} x_i^{\,2} - \frac{1}{n}(\sum_{i=1}^{n} x_i)^2]$$

- **Standard deviation** is the square root of variance s^2

Visualization

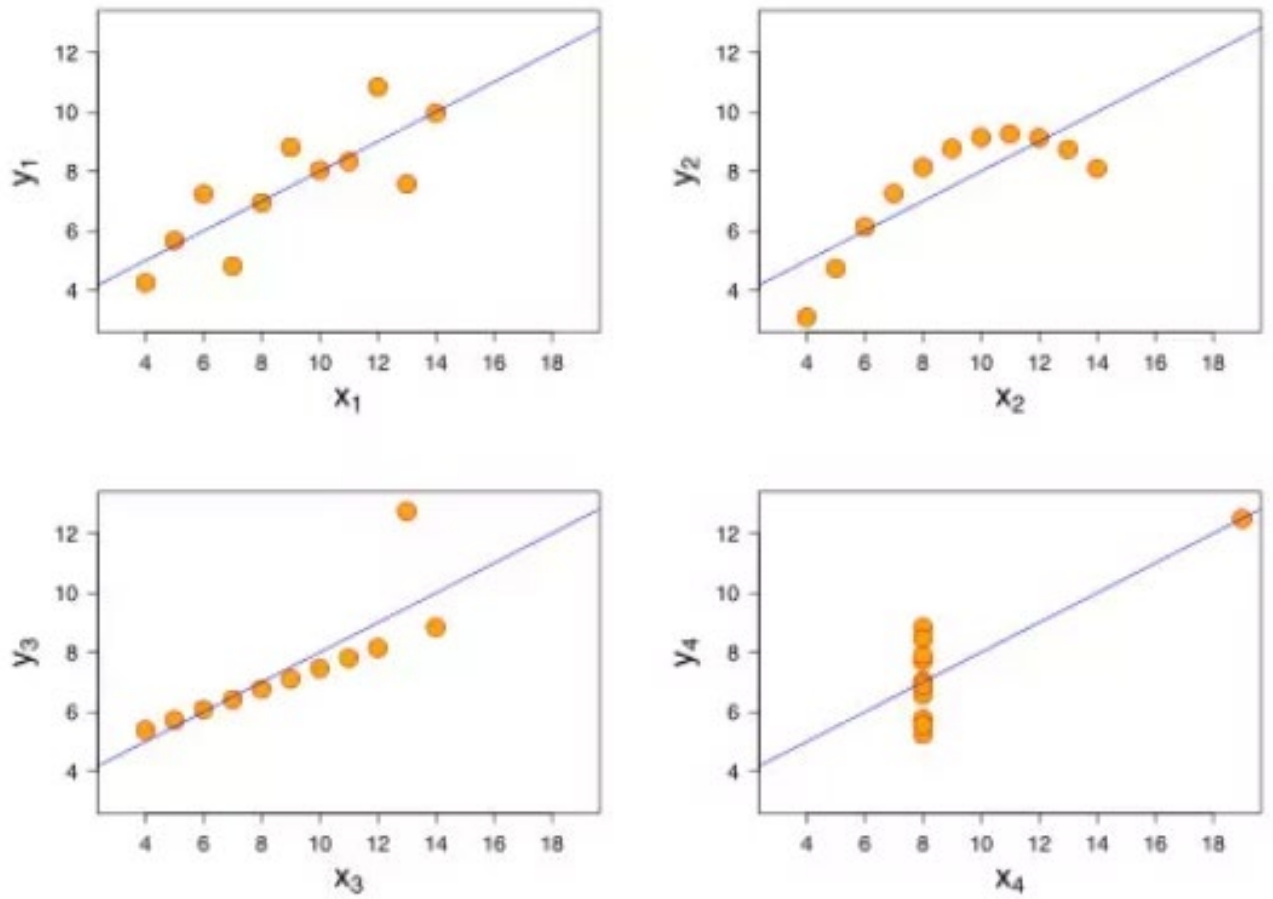

Visualization

- Box-plot
- Histogram
- Scatter-plot
- Quantile and Quantile-Quantile Plot

Boxplot

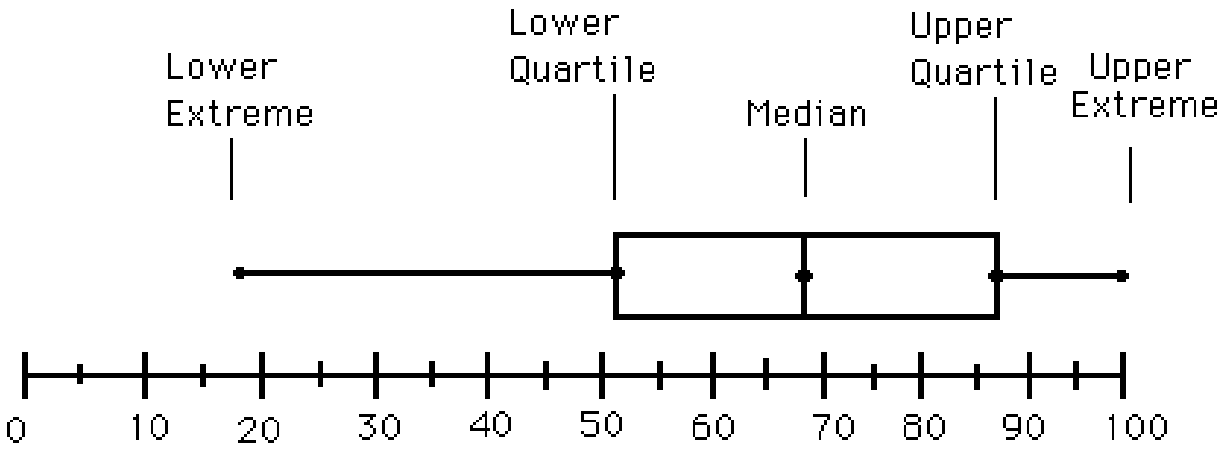

- **Five-number summary** of a distribution
  - Minimum, Q1, Median, Q3, Maximum
- **Boxplot**
  - Data is represented with a box
  - The ends of the box are at the first and third quartiles, i.e., the height of the box is IQR
  - The median is marked by a line within the box
  - Whiskers: two lines outside the box extended to Minimum and Maximum
  - Outliers: points beyond a specified outlier threshold, plotted individually

Histogram

- Histogram: Graph display of tabulated frequencies, shown as bars
- It shows what proportion of cases fall into each of several categories
- Differs from a bar chart in that it is the *area* of the bar that denotes the value, not the height as in bar charts, a crucial distinction when the categories are not of uniform width
- The categories are usually specified as non-overlapping intervals of some variable. The categories (bars) must be adjacent

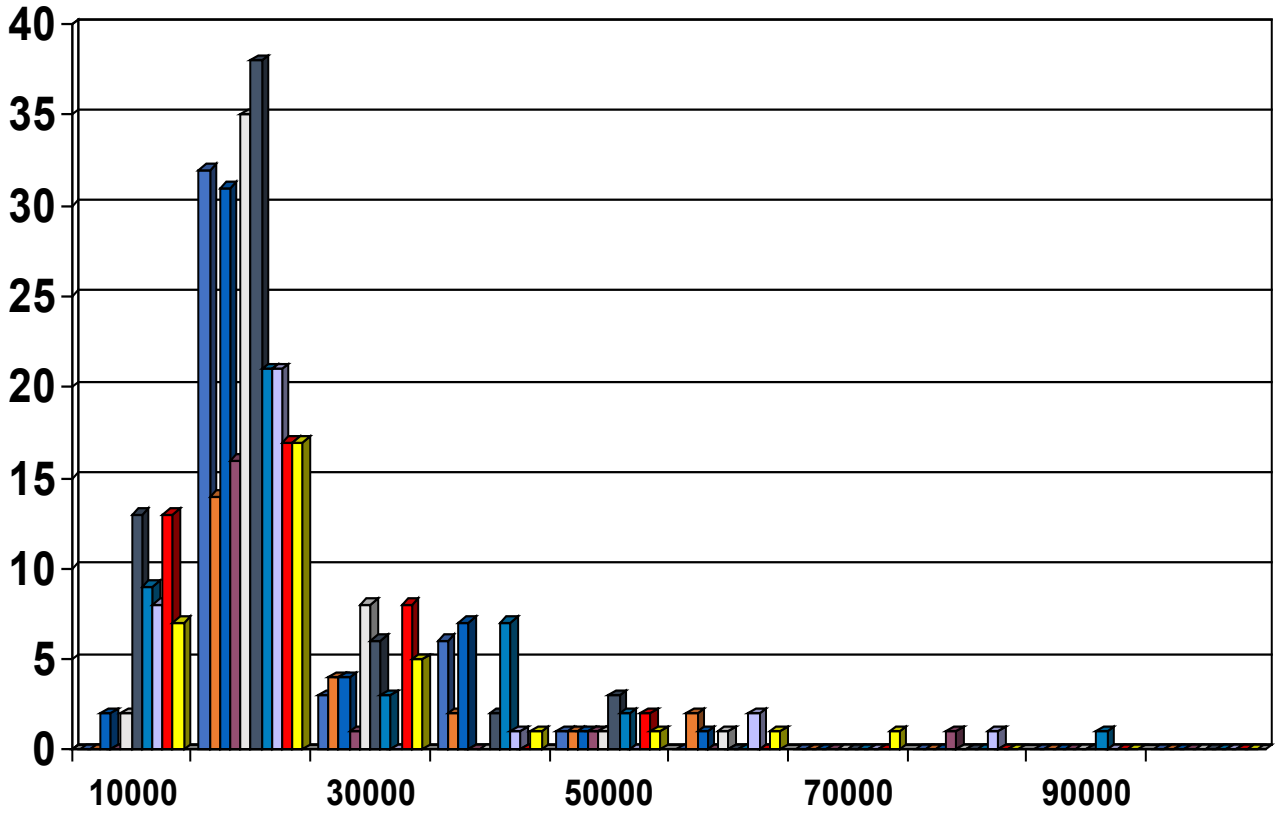

Histograms Often Tell More than Boxplots

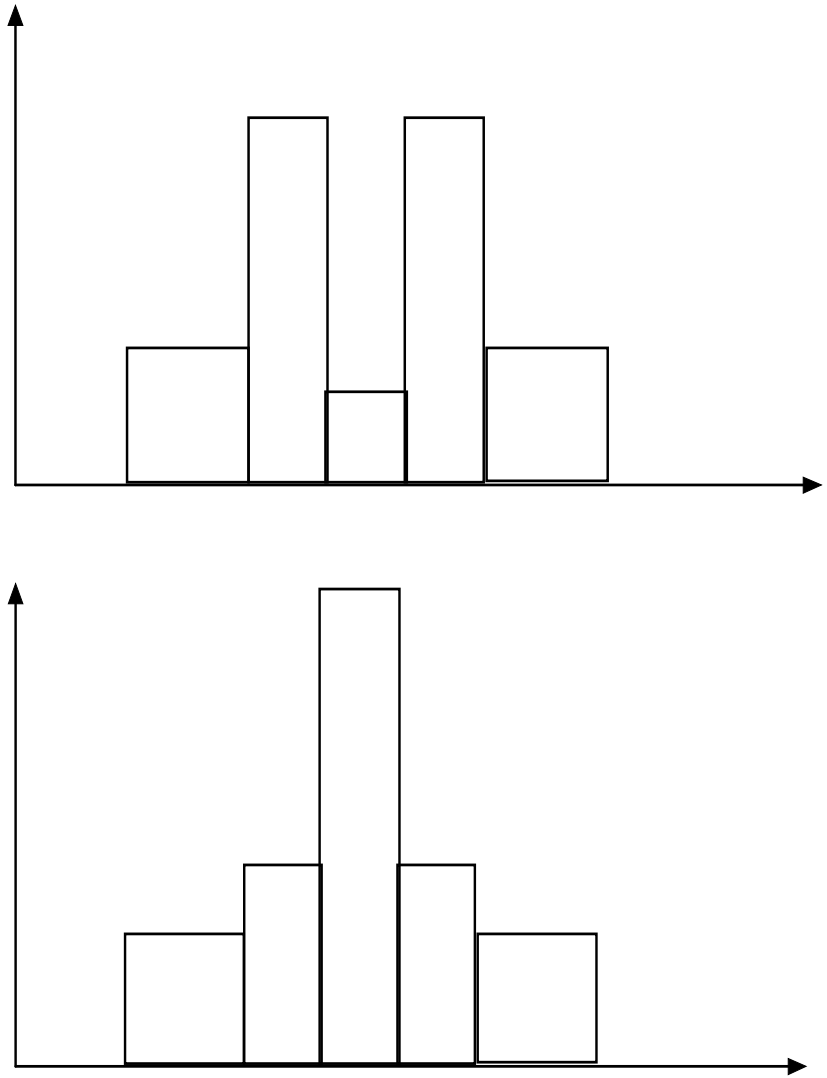

- The two histograms shown in the left may have the same boxplot representation
  - The same values for: min, Q1, median, Q3, max
- But they have rather different data distributions

Quantile Plot

- Plots **quantile** information
  - For a dataset consisting of points x_i sorted in increasing order, f_i indicates that approximately 100 f_i% of the data are below or equal to the value x_i

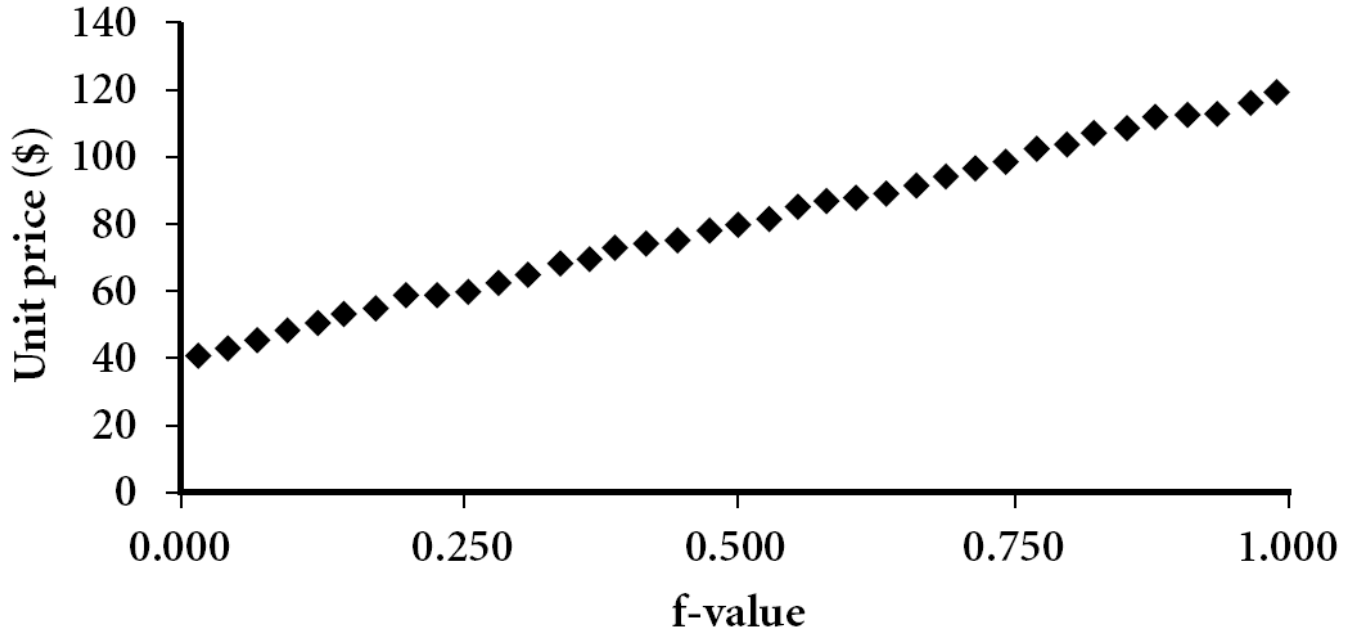

Quantile-Quantile (Q-Q) Plot

- Graphs the quantiles of one univariate distribution against the corresponding quantiles of another
  - Example shows unit price of items sold at Branch 1 vs. Branch 2 for each quantile. Unit prices of items sold at Branch 1 tend to be lower than those at Branch 2.

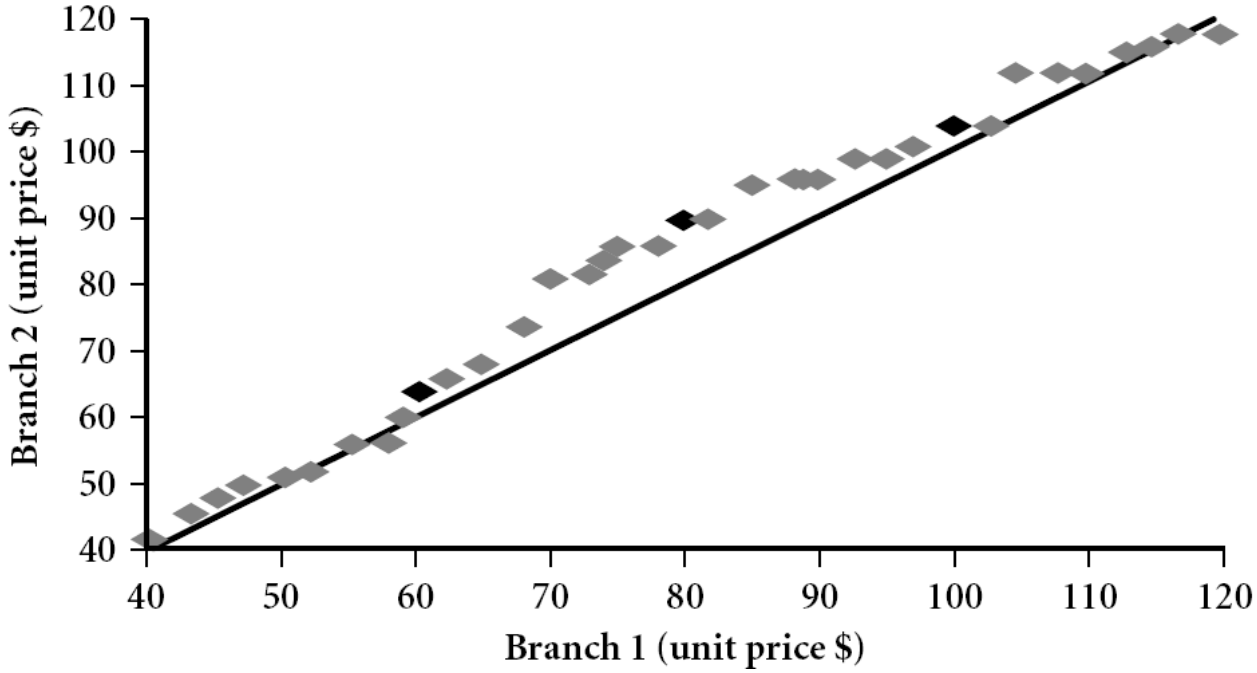

Scatter plot

- Provides a first look at bivariate data to see clusters of points, outliers, etc
- Each pair of values is treated as a pair of coordinates and plotted as points in the plane

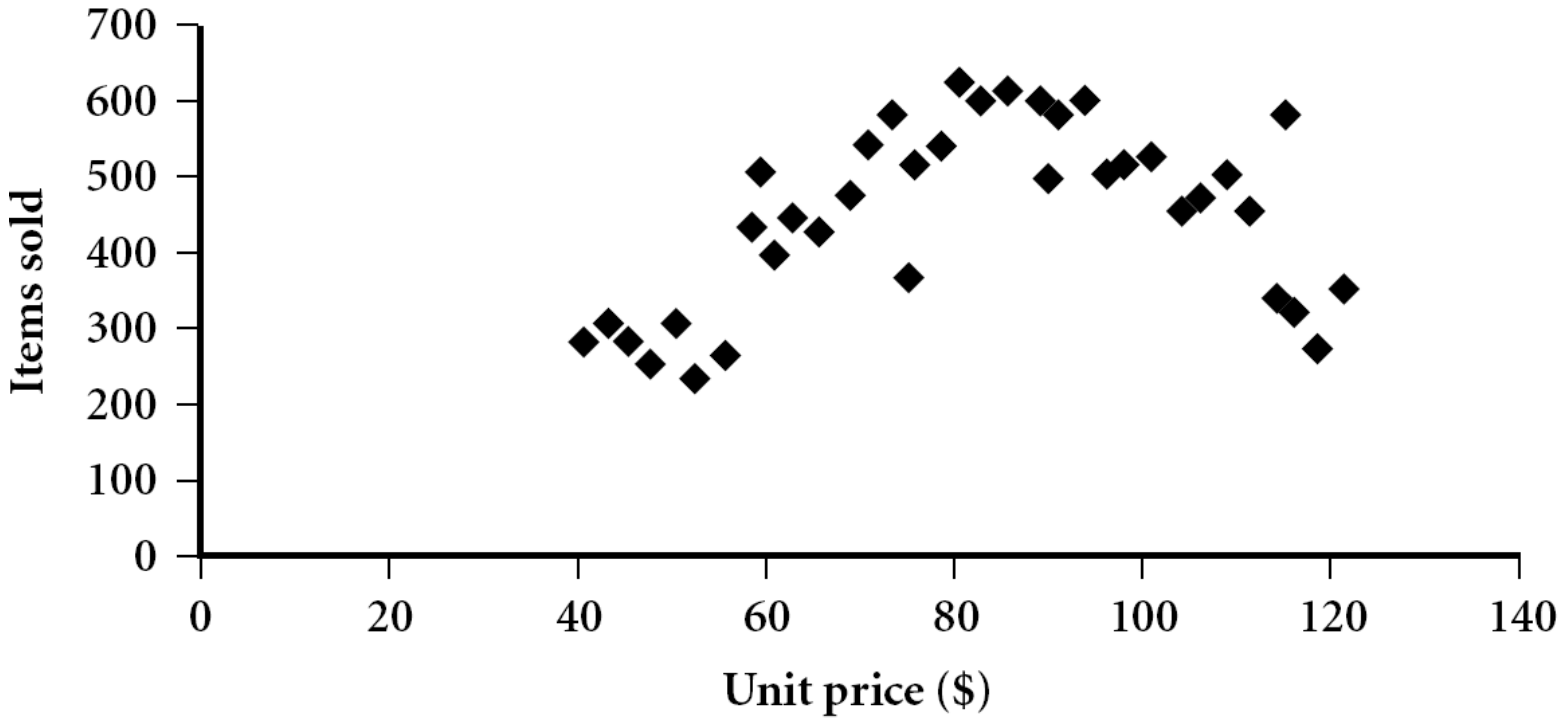

Positively and Negatively Correlated Data

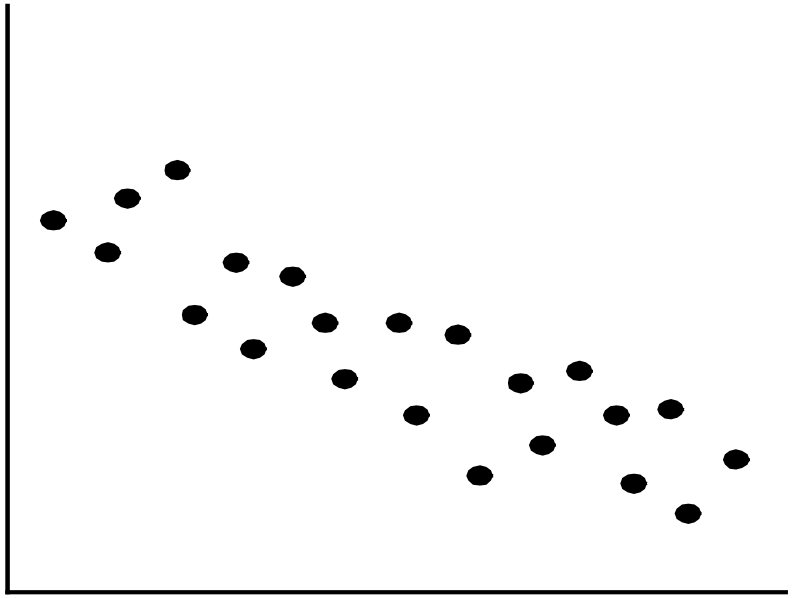

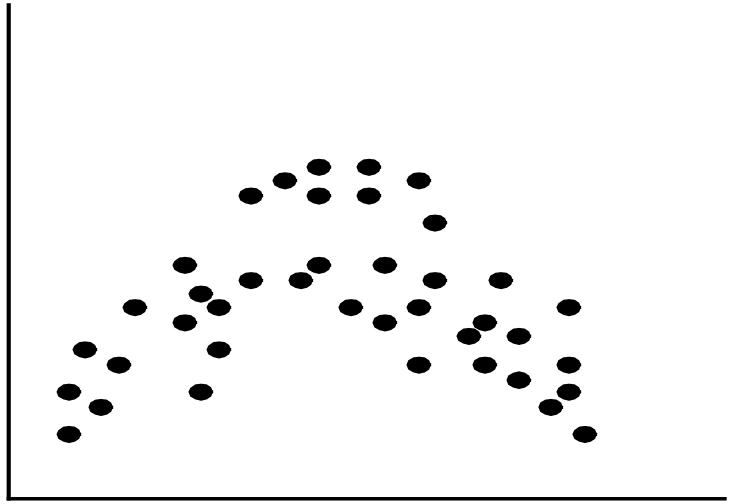

- The left half fragment is positively correlated
- The right half is negative correlated

Visualizations

- **Boxplot**: graphic display of five-number summary
- **Histogram**: x-axis are values, y-axis repres. frequencies
- **Quantile plot**: each value x_i is paired with f_i indicating that approximately 100 f_i % of data are $\leq x_i$
- **Quantile-quantile (q-q) plot**: graphs the quantiles of one univariant distribution against the corresponding quantiles of another
- **Scatter plot**: each pair of values is a pair of coordinates and plotted as points in the plane

Summary

- Data engineering involves knowledge discovery and data mining
- Business knowledge is important for adoption of technologies supporting data mining.
- Properties of qualitative and quantitative data
- Numerical summaries of data such as mean, median, mode
- Visualization of data including histogram and scatter-plot