


Data Engineering

MG-GY 8441

Processing Data

Data cleaning




Data transformation


$-2, 32, 100, 59, 48 \longrightarrow -0.02, 0.32, 1.00, 0.59, 0.48$

Processing Data

Data integration



Processing Data



Storing Data

- Agenda
 - Data Formats
 - Database Management Systems
 - Online Analytical Processing
 - Structured Query Language
- References
 - Han, Kamber, Pei, *Data Mining: Concepts and Techniques* (Chapter 4.1 – 4.3)
 - (Optional) Garcia-Molina, Ullman, Widom, *Database Systems: The Complete Book* (Chapter 2)

Example

Example from Commerce

Suppose you are a business analyst within the accounting group of a retailer.

Your group has been studying sales from movie streaming. You want to find a competitive advantage through offering popular movies.

You decide to integrate data from your sales records and the IMDB database.

Example

nconst	primaryName	birthYear	deathYear	primaryProfession	knownForTitles
nm0000001	Fred Astaire	1899	1987	soundtrack,actor,miscellaneous	tt0043044,tt0053137,tt0072308,tt0050419
nm0000002	Lauren Bacall	1924	2014	actress,soundtrack	tt0117057,tt0037382,tt0038355,tt0071877
nm0000003	Brigitte Bardot	1934	\N	actress,soundtrack,producer	tt0057345,tt0059956,tt0049189,tt0054452
nm0000004	John Belushi	1949	1982	actor,soundtrack,writer	tt0072562,tt0080455,tt0078723,tt0077975
nm0000005	Ingmar Bergman	1918	2007	writer,director,actor	tt0083922,tt0060827,tt0050976,tt0050986

Example


Example from Commerce

You want to compare the revenue generated by a movie and the content of the movie


- How does genre impact sales?
- Do some directors more popular?
- Will customers respond to indications of ratings?

Data Formats


Structured Data



Semi-Structured Data



Unstructured Data



Unstructured Data

```
Log: Log file open, 06/10/18 16:28:00
Log: WinSock: version 1.1 (2.2), MaxSocks=32767, MaxUdp=65467
Log: Version: 8630
Log: Compiled (32-bit): Sep  3 2015 21:05:18
Log: Changelist: 1100103
Log: Command line:
```

Semi-Structured Data

XML	JSON	YAML
<pre><Servers> <Server> <name>Server1</name> <owner>John</owner> <created>123456</created> <status>active</status> </Server> </Servers></pre>	<pre>{ Servers: [{ name: Server1, owner: John, created: 123456, status: active }] }</pre>	<pre>Servers: - name: Server1 owner: John created: 123456 status: active</pre>

Structured Data

CSV

```
Candidate,Party,%,Year,Result
Reagan,Republican,50.7,1980,win
Carter,Democratic,41,1980,loss
Anderson,Independent,6.6,1980,loss
Reagan,Republican,58.8,1984,win
Mondale,Democratic,37.6,1984,loss
Bush,Republican,53.4,1988,win
Dukakis,Democratic,45.6,1988,loss
Clinton,Democratic,43,1992,win
Bush,Republican,37.4,1992,loss
Perot,Independent,18.9,1992,loss
Clinton,Democratic,49.2,1996,win
Dole,Republican,40.7,1996,loss
Perot,Independent,8.4,1996,loss
Gore,Democratic,48.4,2000,loss
Bush,Republican,47.9,2000,win
Kerry,Democratic,48.3,2004,loss
Bush,Republican,50.7,2004,win
Obama,Democratic,52.9,2008,win
McCain,Republican,45.7,2008,loss
Obama,Democratic,51.1,2012,win
Romney,Republican,47.2,2012,loss
Clinton,Democratic,48.2,2016,loss
Trump,Republican,46.1,2016,win
```

Candidate	Party	%	Year	Result
Reagan	Republican	50.7	1980	win
Carter	Democratic	41.0	1980	loss
Anderson	Independent	6.6	1980	loss
Reagan	Republican	58.8	1984	win
Mondale	Democratic	37.6	1984	loss
Bush	Republican	53.4	1988	win
Dukakis	Democratic	45.6	1988	loss
Clinton	Democratic	43.0	1992	win
Bush	Republican	37.4	1992	loss
Perot	Independent	18.9	1992	loss
Clinton	Democratic	49.2	1996	win
Dole	Republican	40.7	1996	loss
Perot	Independent	8.4	1996	loss
Gore	Democratic	48.4	2000	loss
Bush	Republican	47.9	2000	win
Kerry	Democratic	48.3	2004	loss
Bush	Republican	50.7	2004	win
Obama	Democratic	52.9	2008	win
McCain	Republican	45.7	2008	loss
Obama	Democratic	51.1	2012	win
Romney	Republican	47.2	2012	loss
Clinton	Democratic	48.2	2016	loss
Trump	Republican	46.1	2016	win

tsv

File Systems

- ▶ Computers provide access to a command line interface. Users input commands to perform operations on the computer especially files.

```
!ls
```

```
data  ds-ua-112-lab04.ipynb  movies_100_rows.csv  movies.csv
```

- ▶ We can enter commands in Jupyter notebook using exclamation point
- ▶ Note that the commands differ across operating systems. Here we have the commands for the Linux operating system on JupyterHub.

File Systems

- ▶ Some commands for accessing files include
 - ▶ head
 - ▶ Returns the first 10 rows of the file
 - ▶ tail
 - ▶ Returns the last 10 rows of the file
 - ▶ cat
 - ▶ Returns all rows of the file

```
!head movies.csv
```

```
director,genre,movie,rating,revenue  
David,Action & Adventure,Deadpool 2,7,318344544  
Bill,Comedy,Book Club,5,68566296  
Ron,Science Fiction & Fantasy,Solo: A Star Wars Story,6,213476293  
Baltasar,Drama,Adrift,6,31445012  
Bart,Drama,American Animals,6,2847319  
Gary,Action & Adventure,Oceans 8,6,138803463  
Drew,Action & Adventure,Hotel Artemis,8,6708147  
Brad,Animation,Incredibles 2,5,594398019  
Jeff,Comedy,Tag,6,54336863
```

File Systems

- ▶ A command for determining the size of files is
 - ▶ `du -sh`
 - ▶ Will calculate the size of files or folders

```
!du -sh data
```

```
28K      data
```

```
!du -sh data/*
```


```
12K      data/more_data
```

```
4.0K     data/movies_100_rows.csv
```

```
4.0K     data/movies.csv
```

Database Management Systems





Data Warehouse

Collects and organizes historical data from multiple sources

Data is *periodically* **ETL**ed into the data warehouse:

- **Extracted** from remote sources
- **Transformed** to standard schemas
- **Loaded** into the (typically) relational (SQL) data system

Extract → **T**ransform → **L**oad (ETL)

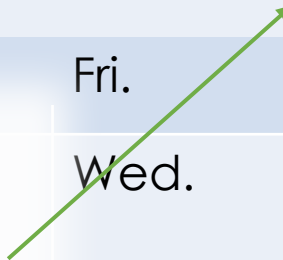
Extract & Load: provides a snapshot of operational data

- Historical snapshot
- Data in a single system
- Isolates analytics queries (e.g., Deep Learning) from business critical services (e.g., processing user purchase)
- Easy!

Transform: clean and prepare data for analytics in a unified representation

- **Difficult** → often requires specialized code and tools
- Different schemas, encodings, granularities

Table

pname	category	price	qty	date	day	city	state	country
Corn	Food	25	25	3/30/16	Wed.	Omaha	NE	USA
Corn	Food	25	8	3/31/16	Thu.	Omaha	NE	USA
Corn	Food	25	15	4/1/16	Fri.	Omaha	NE	USA
 ➤ Substantial redundancy → expensive to store					Wed.	Omaha	NE	USA
					Thu.	Omaha	NE	USA
					Fri.	Omaha	NE	USA
Galaxy 1	Phones	18	8	1/30/16	Wed.	Omaha	NE	USA
Peanuts	Food	2	45	3/31/16	Thu.	Seoul		Korea

Multidimensional Data Model

Sales **Fact Table**

pid	timeid	locid	sales
11	1	1	25
11	2	1	8
11	3	1	15
12	1	1	30
12	2	1	20
12	3	1	50
12	1	1	8
13	2	1	10
13	3	1	10
11	1	2	35
11	2	2	22
11	3	2	10
12	1	2	26

Locations

locid	city	state	country
1	Omaha	Nebraska	USA
2	Seoul		Korea
5	Richmond	Virginia	USA

Products


pid	pname	category	price
11	Corn	Food	25
12	Galaxy 1	Phones	18
13	Peanuts	Food	2

Time

timeid	Date	Day
1	3/30/16	Wed.
2	3/31/16	Thu.
3	4/1/16	Fri.

Dimension Tables

➤ Multidimensional “Cube” of data



Multidimensional Data Model

Sales **Fact Table**

pid	timeid	locid	sales
11	1	1	25
11	2	1	8
11	3	1	15
12	1	1	30
12	2	1	20
12	3	1	50
12	1	1	8
13	2	1	10
13	3	1	10
11	1	2	35
11	2	2	22
11	3	2	10
12	1	2	26

Locations

locid	city	state	country
1	Omaha	Nebraska	USA
2	Seoul		Korea
5	Richmond	Virginia	USA

Products

pid	pname	category	price
11	Corn	Food	25
12	Galaxy 1	Phones	18
13	Peanuts	Food	2


Time

timeid	Date	Day
1	3/30/16	Wed.
2	3/31/16	Thu.
3	4/1/16	Fri.

Dimension Tables

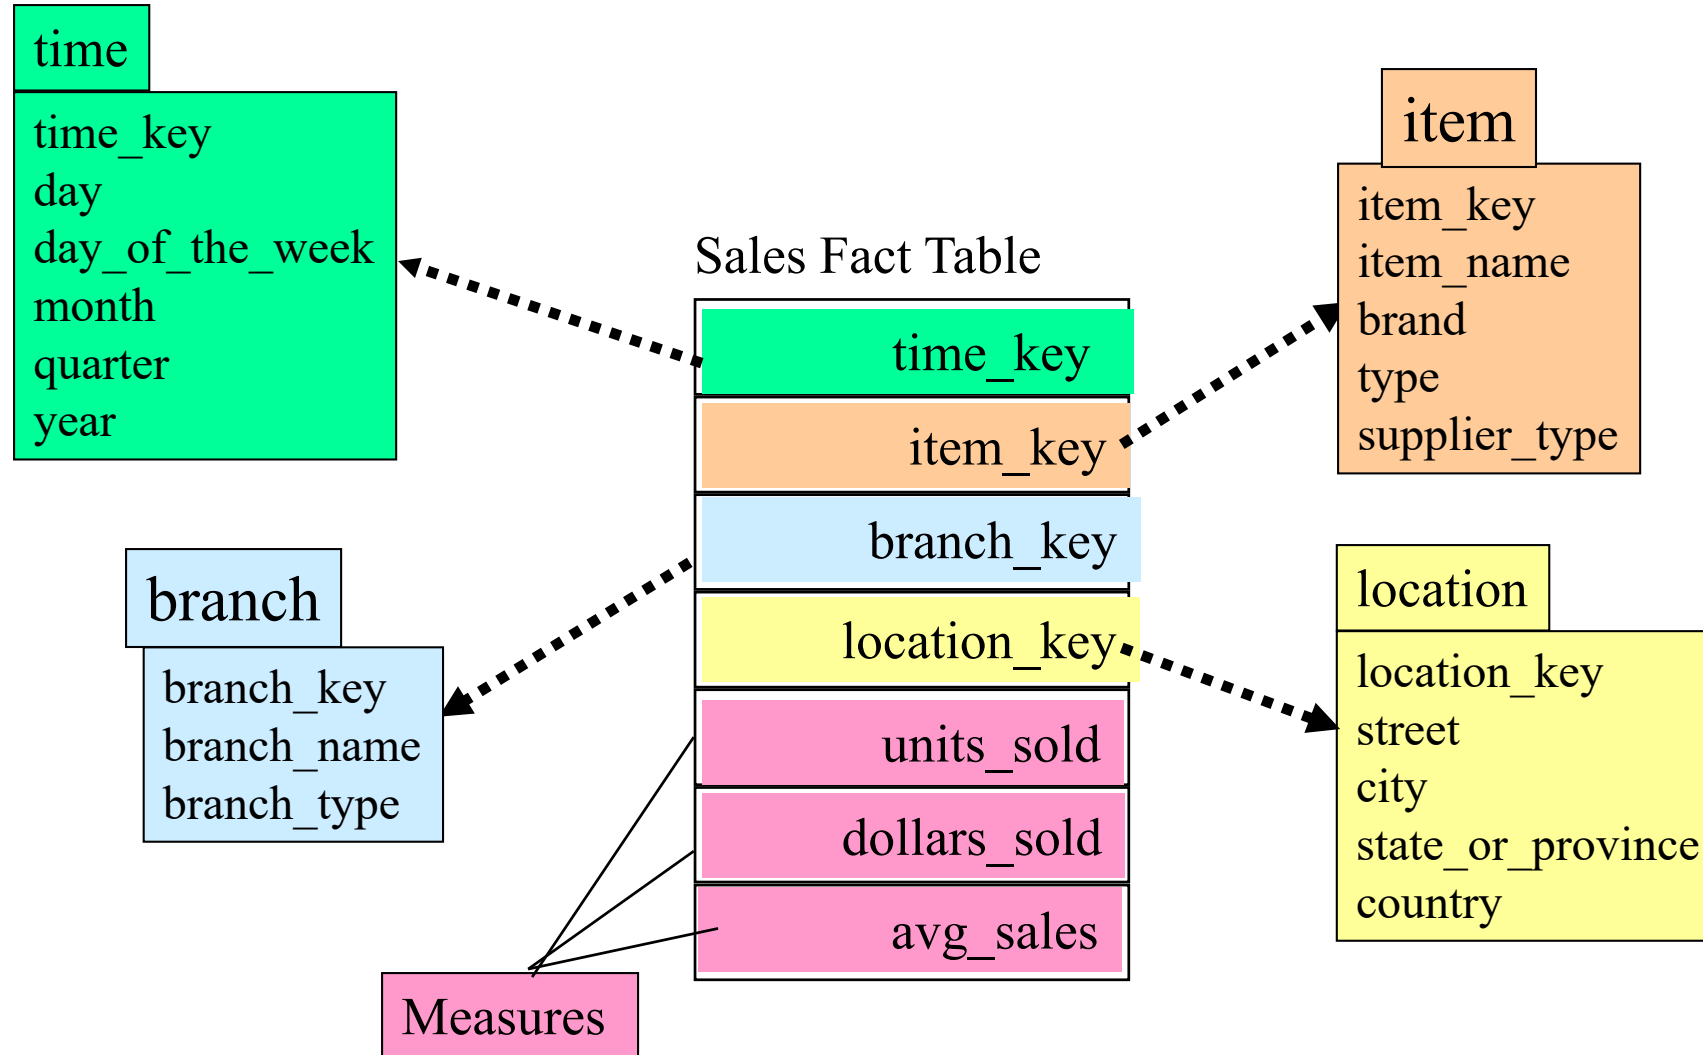
- Fact Table
 - minimizes redundant info.
 - Reduces data errors
- Dimensions
 - easy to manage, summarize and rename
- How can we combine tables?
 - Joins

The Star Schema




← This looks like a star ...

Star Schema: An Example

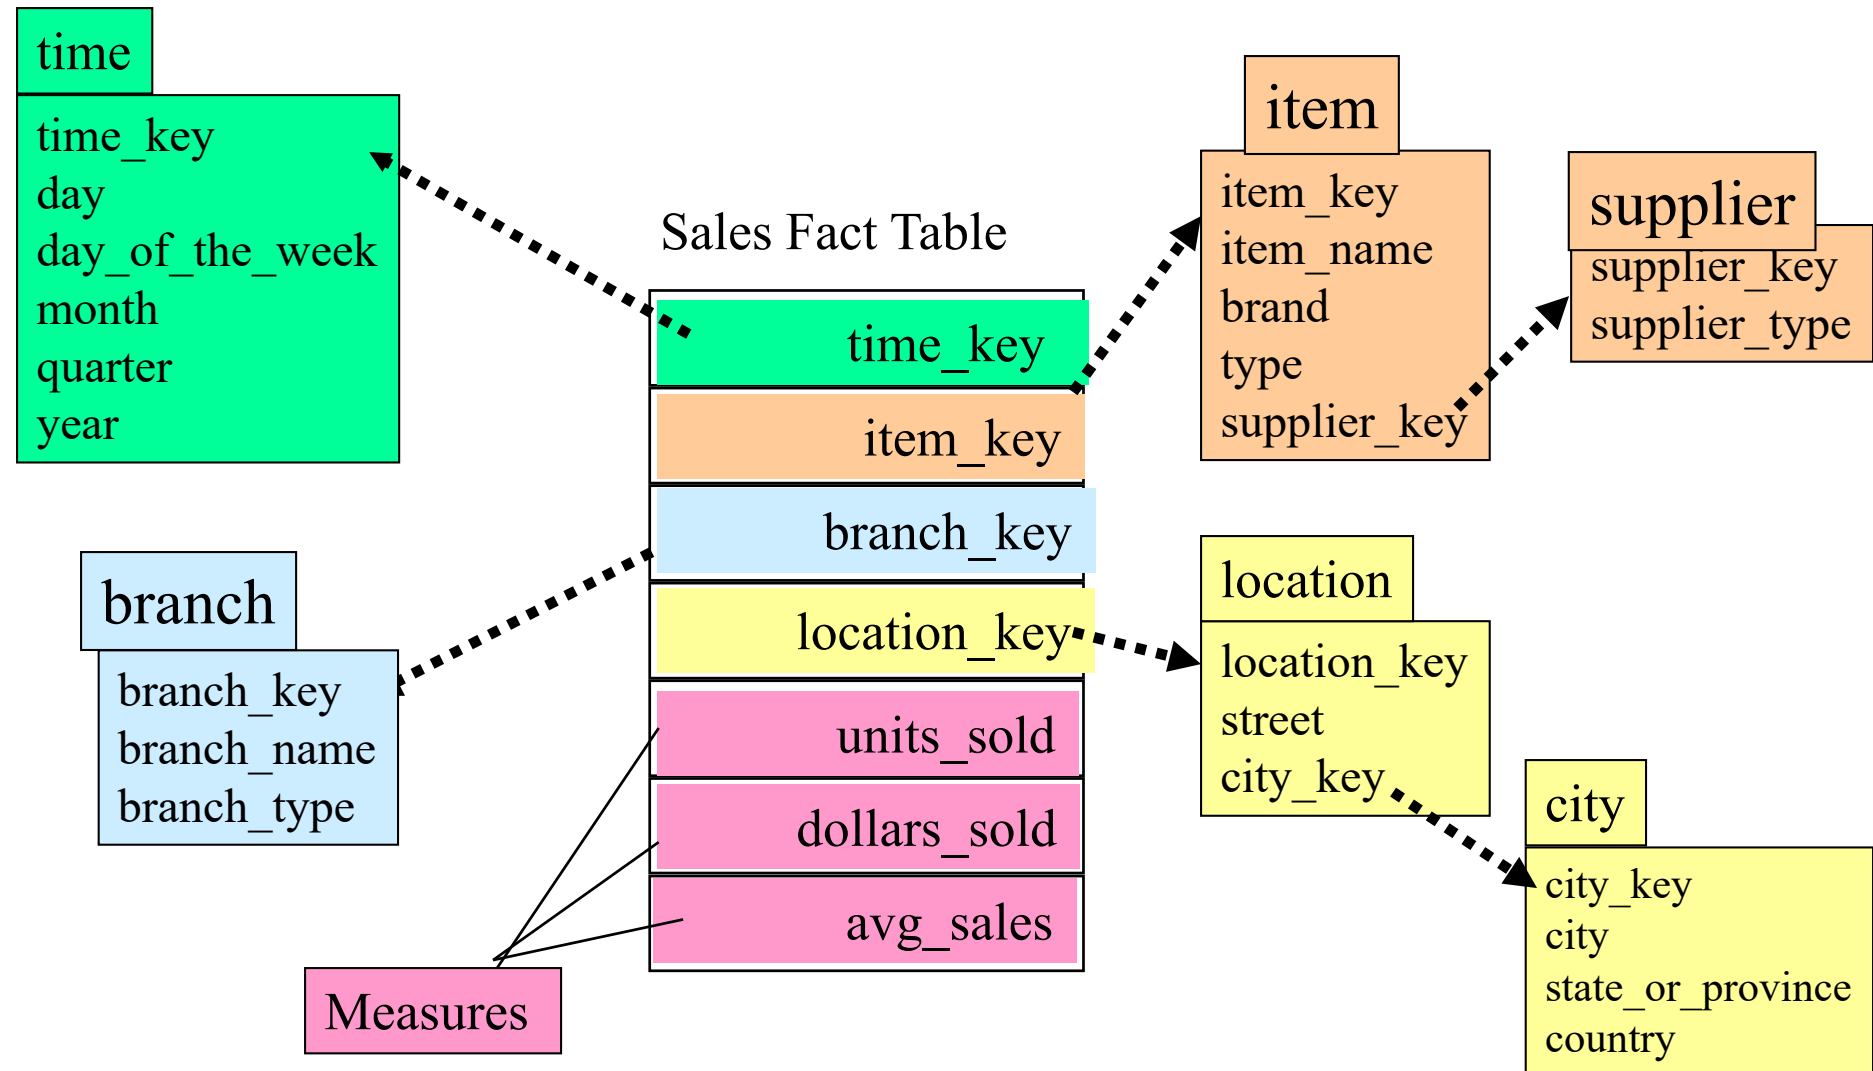


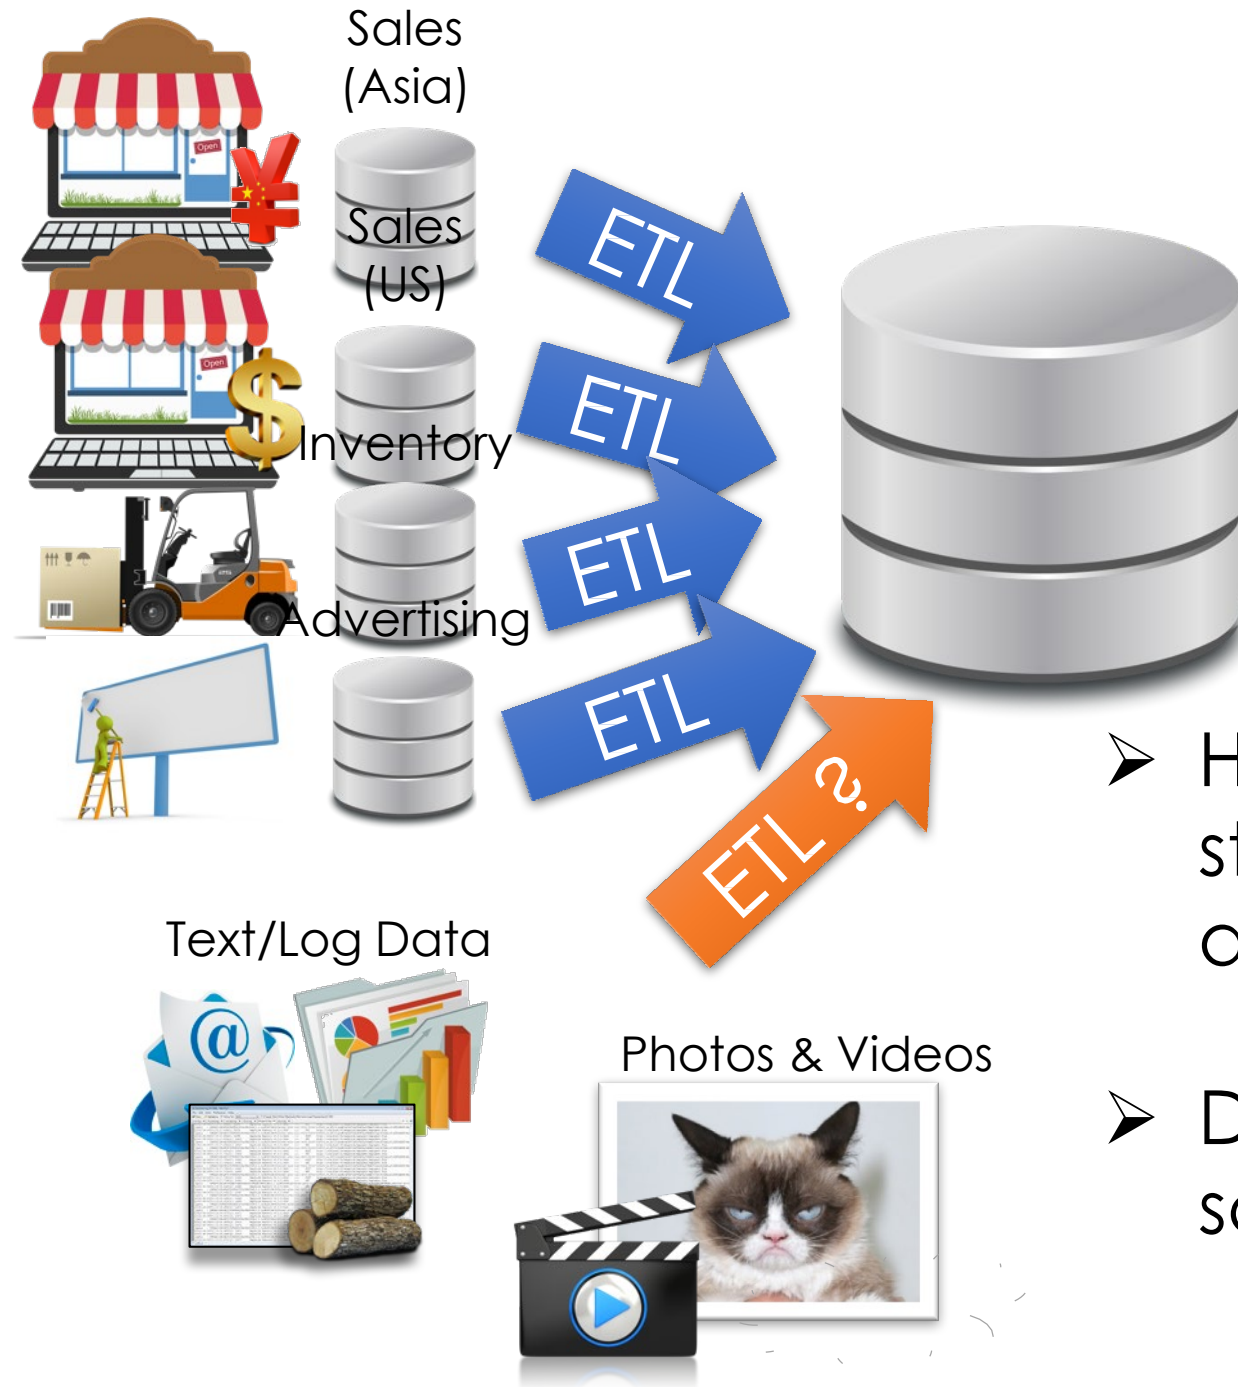
The Snowflake Schema



← This looks like a snowflake
...sort of

Snowflake Schema: An Example





Data Warehouse

Collects and organizes historical data from multiple sources

- How do we deal with semi-structured and unstructured data?
- Do we really want to force a schema on load?

Sales
(Asia)

How do we **clean**
and **organize** this
data?

Depends on use ...



Data Warehouse

Collects and organizes
historical data from
multiple sources



Text/Log Data



Photos & Videos



How do we **load** and **process** this
data in a relational system?

Depends on use ...
Can be difficult ...
Requires thought ...



Data Lake

Store a copy of all the data

- in one place
- in its original “natural” form

Enable data consumers to choose how to transform and use data.

- *Schema on Read*

Enabled by new Tools:

Map-Reduce & Distributed Filesystems

What could go wrong?

It is Terrible!

Online Analytical Processing



OLTP vs. OLAP

❑ OLTP: Online transactional processing

❑ DBMS operations

❑ Query and transactional processing


❑ OLAP: Online analytical processing

❑ Data warehouse operations

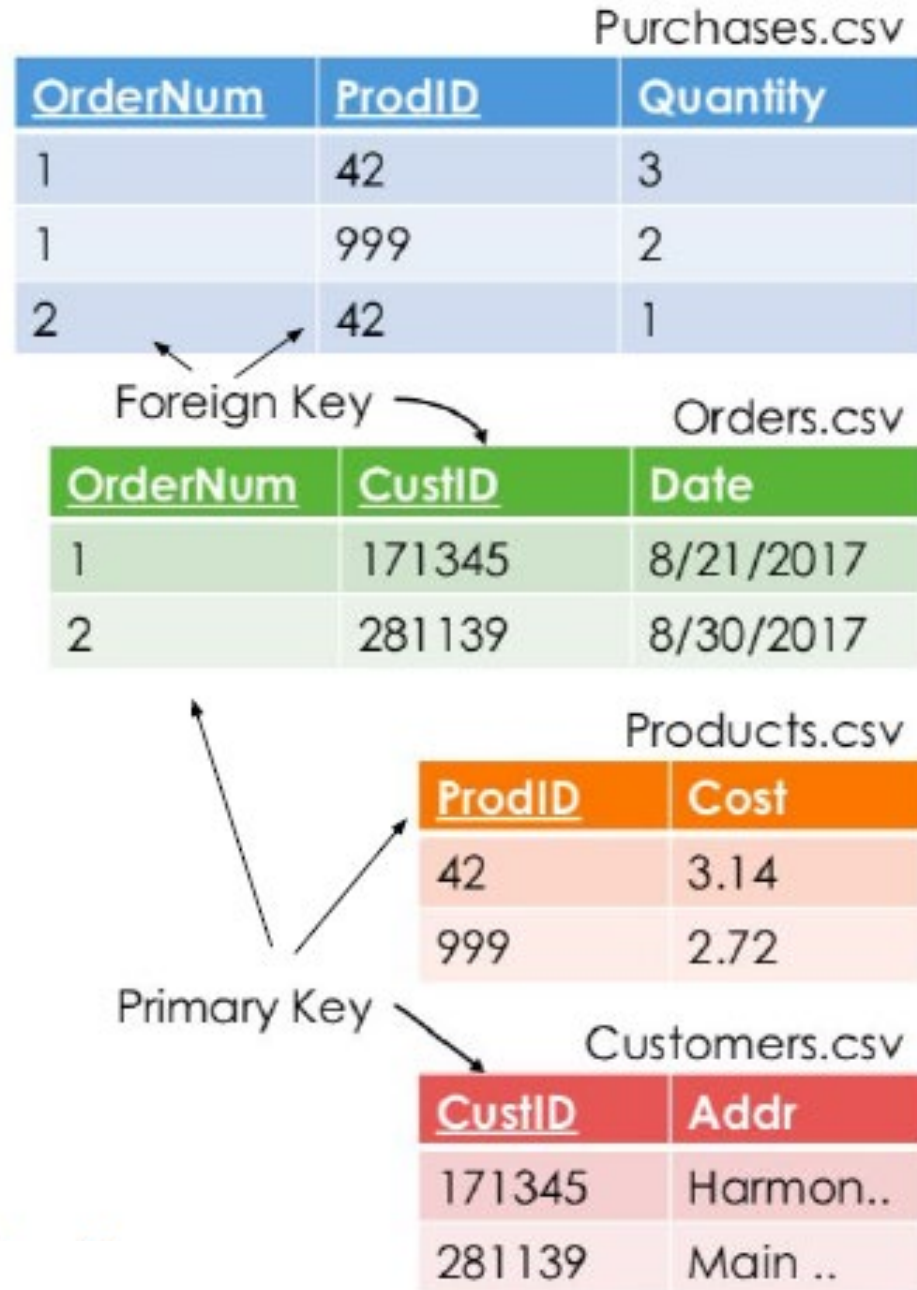
❑ Drilling, slicing, dicing, etc.

	OLTP	OLAP
users	clerk, IT professional	knowledge worker
function	day to day operations	decision support
DB design	application-oriented	subject-oriented
data	current, up-to-date detailed, flat relational isolated	historical, summarized, multidimensional integrated, consolidated
usage	repetitive	ad-hoc
access	read/write index/hash on prim. key	lots of scans
unit of work	short, simple transaction	complex query
# records accessed	tens	millions
#users	thousands	hundreds
DB size	100MB-GB	100GB-TB
metric	transaction throughput	query throughput, response

Data Cube



Keys



Index

Base table

Cust	Region	Type
C1	Asia	Retail
C2	Europe	Dealer
C3	Asia	Dealer
C4	America	Retail
C5	Europe	Dealer

Index on Region

RecID	Asia	Europe	America
1	1	0	0
2	0	1	0
3	1	0	0
4	0	0	1
5	0	1	0


Index on Type

RecID	Retail	Dealer
1	1	0
2	0	1
3	0	1
4	1	0
5	0	1

ACID

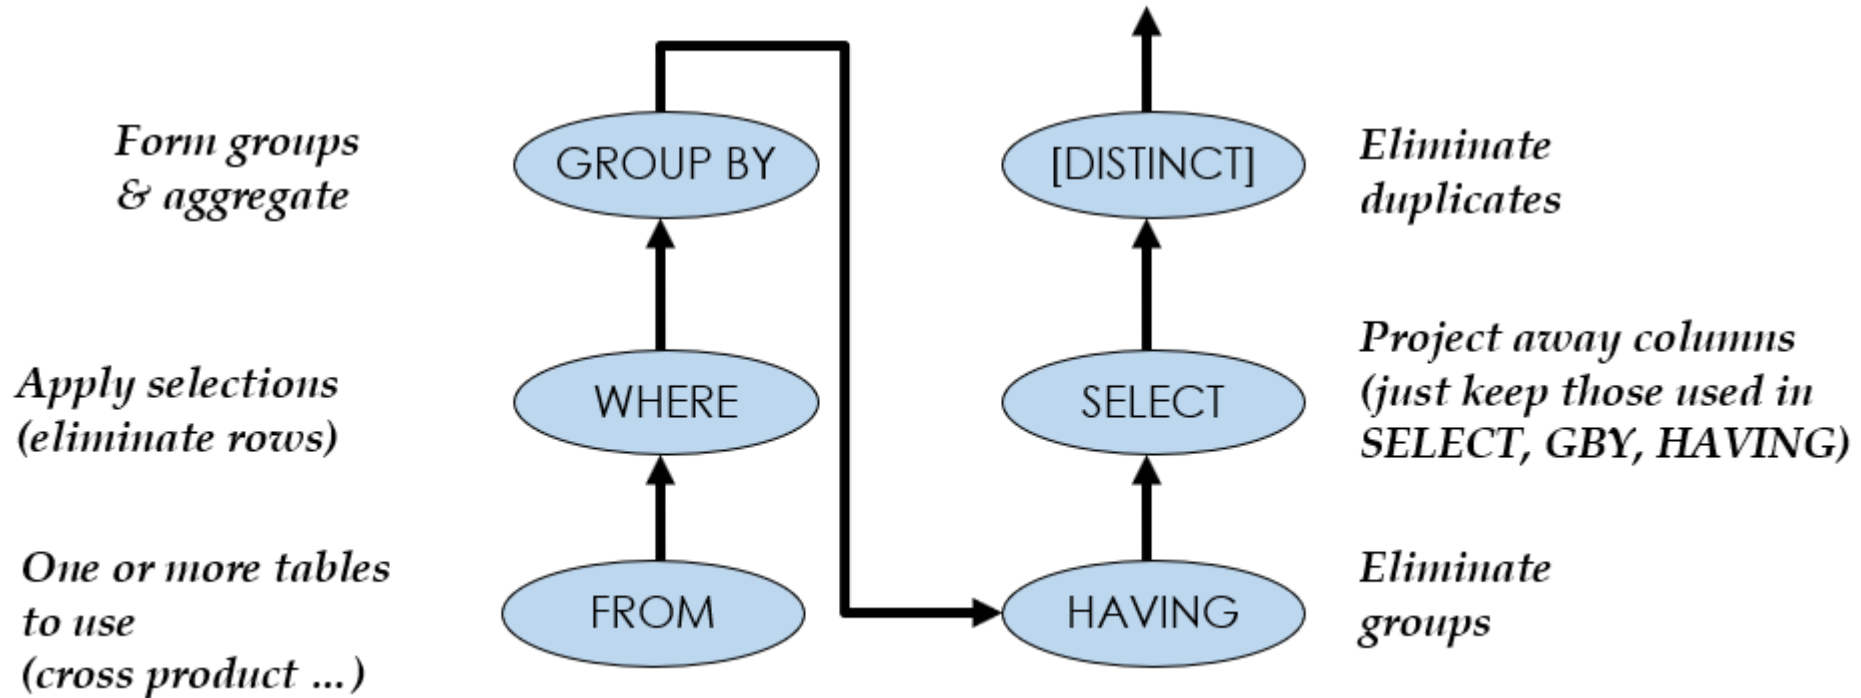
Atomicity	Operations are all-or-nothing (No partial updates; operations bundled in transactions)
Consistency	Transactions move from one valid state to another
Isolation	Concurrent operations do not depend on order of execution
Durability	Completed transactions are permanent (usually implemented by flushing to disk before completion)

Structured Query Language



Commands

SELECT	[DISTINCT] target-list
FROM	relation-list
WHERE	qualification
GROUP BY	grouping-list
HAVING	group-qualification




Cross Join

- ▶ Cross Join pairs each of the rows in the tables regardless of the entries in the columns.
- ▶ All pairs of rows appear in the Join.

s	
<u>M</u>	<u>U</u>
1	W
2	X
3	X
4	Y

t	
<u>N</u>	<u>V</u>
A	X
B	X
C	Y
D	Z


<u>M</u>	<u>U</u>	<u>N</u>	<u>V</u>
1	W	A	X
2	X	A	X
3	X	A	X
4	Y	A	X
1	W	B	X
2	X	B	X
3	X	B	X
4	Y	B	X



Inner Join

- ▶ Inner Join pairs each of the rows in the tables depending on the entries in specific columns.
- ▶ The entries of columns must match for the pair to appear in the Join.

INNER JOIN




s		t	
<u>M</u>	<u>U</u>	<u>N</u>	<u>V</u>
1	W	A	X
2	X	A	X
3	X	B	X
4	Y	C	Y
		D	Z

<u>M</u>	<u>U</u>	<u>N</u>	<u>V</u>
2	X	A	X
3	X	A	X
2	X	B	X
3	X	B	X
4	Y	C	Y

Left Join

- ▶ Left Join pairs each of the rows in the left table to rows in the right table depending on the entries in specific columns.
- ▶ The entries of columns in the right table must match for the pair to appear in the Join.

LEFT JOIN




s		t	
<u>M</u>	<u>U</u>	<u>N</u>	<u>V</u>
1	W		
2	X	A	X
3	X	B	X
4	Y	C	Y
		D	Z

1	W	null	null
2	X	A	X
3	X	A	X
2	X	B	X
3	X	B	X
4	Y	C	Y

Right Join

- ▶ Right Join pairs each of the rows in the right table to rows in the left table depending on the entries in specific columns.
- ▶ The entries of columns in the left table must match for the pair to appear in the Join.



s		t	
<u>M</u>	<u>U</u>	<u>N</u>	<u>V</u>
1	W	A	X
2	X	B	X
3	X	C	Y
4	Y	D	Z

2	X	A	X
3	X	A	X
2	X	B	X
3	X	B	X
4	Y	C	Y
null	null	D	Z

Outer Join


- ▶ Outer Join combines Left Join and Right Join
- ▶ Note that Outer Join does not contain duplicate entries for the matching rows, that is, the rows contained in the Inner Join.

s	
<u>M</u>	<u>U</u>
1	W
2	X
3	X
4	Y

t	
<u>N</u>	<u>V</u>
A	X
B	X
C	Y
D	Z

1	W	null	null
2	X	A	X
3	X	A	X
2	X	B	X
3	X	B	X
4	Y	C	Y
null	null	D	Z

FULL OUTER JOIN



Summary

- Data Formats
 - structured, unstructured, semi-structured
- Database Management Systems
 - extract, transform, load
- Online Analytical Processing
 - data cube, keys and indices, slicing/drilling
- Structured Query Language
 - joins